

F5 Training

# Securing Applications with NGINX

Lab Only Guide



NGINX Plus R29

Publishing Version: 2023-06-16



## Lab Only Guide

**2023-06-16**

## Support and Contact Information

### Obtaining Technical Support

**Web:** [support.f5.com](https://support.f5.com) (Ask F5)

**Suggestions:** [feedback@f5.com](mailto:feedback@f5.com)



### Contacting F5

**Web:** [www.f5.com](https://www.f5.com)

**Sales:** [sales@f5.com](mailto:sales@f5.com)

**Information:** [info@f5.com](mailto:info@f5.com)



### Contact Training

**F5 Information:** [www.f5.com/services/training](https://www.f5.com/services/training)

**North America/South America:** [training@f5.com](mailto:training@f5.com)

**Europe/Middle East/Africa:** [emeatraining@f5.com](mailto:emeatraining@f5.com)

**Japan:** [japantraining@f5.com](mailto:japantraining@f5.com)

**Asia/Pacific:** [apactraining@f5.com](mailto:apactraining@f5.com)

**NGINX Information:** [www.nginx.com/learn/nginx-training](https://www.nginx.com/learn/nginx-training)

**NGINX Training:** [NGINX-training@f5.com](mailto:NGINX-training@f5.com)



## Legal Notices

Copyright ©2022 F5, Inc. All rights reserved.

F5, Inc. (F5) believes the information it furnishes to be accurate and reliable. However, F5 assumes no responsibility for the use of this information; nor any infringement of patents or other rights of third parties which may result from its use. No license is granted by implication or otherwise under any patent; copyright; or other intellectual property right of F5 except as specifically described by applicable user licenses. F5 reserves the right to change specifications at any time without notice.

## Trademarks

Access Policy Manager, APM, Advanced Firewall Manager, AFM, Advanced Routing, Advanced Web Application Firewall, Advanced WAF, Adv. WAF, Application Security Manager, ASM, AskF5, BIG-IP, BIG-IQ, Clustered Multiprocessing, CMP, DevCentral, DNS Express, Edge Client, Edge Gateway, Edge Portal, F5, F5 Agility, F5OS, Fast Application Proxy, Fast Cache, Global Traffic Manager, GTM, iApps, iControl, iHealth, iQuery, iRules, Local Traffic Manager, LTM, NGINX, NGINX Plus, NGINX Controller, NGINX Service Mesh, NGINX App Protect, NGINX Instance Manager, NGINX Unit, NGINX Amplify, NGINX Kubernetes Ingress Controller, OneConnect, Policy Enforcement Manager, PEM, Protocol Security Manager, PSM, Silverline, SSL Acceleration, SSL Everywhere, SSL Orchestrator, SYN Check, TCP Express, Traffic Management Operating System, TMOS, Virtual Clustered Multiprocessing, vCMP, VELOS, and VIPRION are trademarks or service marks of F5, Inc., in the U.S. and other countries, and may not be used without F5's express written consent. All other product and company names herein may be trademarks of their respective owners.

For more information on F5, Inc. Trademarks, please visit: [www.f5.com/company/policies/trademarks](https://www.f5.com/company/policies/trademarks)



## Materials and Patents

The material reproduced on this manual, including but not limited to graphics, text, pictures, photographs, layout, and the like ("Content"), are protected by United States Copyright law. Absolutely no Content from this manual may be copied, reproduced, exchanged, published, sold, or distributed without the prior written consent of F5, Inc. The information in this document has been carefully verified and is believed to be accurate. F5 assumes no responsibilities for any inaccuracies that may appear in this document. In no event will F5 be liable for direct, indirect, special, exemplary, incidental, or consequential damages resulting from any defect or omission in this document, even if advised of the possibility of such damages.

This product may be protected by one or more patents indicated at: [www.f5.com/company/policies/patents](https://www.f5.com/company/policies/patents)



# Table of Contents

<b>Support and Contact Information .....</b>	<b>i</b>
<b>Legal Notices.....</b>	<b>i</b>
<b>Chapter 1 .....</b>	<b>1</b>
Labs - Introduction to Lab Environment.....	1
Lab 1: Login to Your Training Machine.....	3
<b>Chapter 2 .....</b>	<b>7</b>
Labs - Setting up an HTTPS Server .....	7
Lab 1: Generating Certs and Keys .....	9
Lab 2: Configuring an HTTPS Server .....	12
Lab 3: Set SSL Ciphers .....	16
Lab 4: Configure Forward Secrecy .....	19
Lab 5: Testing HTTPS Configuration.....	21
Labs - Secure Upstream Traffic.....	25
Lab 1: Secure Upstream Traffic.....	26
Lab 2: Setup the NGINX Plus API Dashboard .....	29
Lab 3: Configure NGINX Logging .....	33
<b>Chapter 3 .....</b>	<b>37</b>
Labs - Set Limit Rates.....	37
Lab 1: Configure NGINX limit rates .....	38
Labs - Logs .....	43
Lab 1: NGINX Access and Error Log Files .....	44
Lab 2: Logging Limit Rates .....	49
Labs - Setup Dynamic IP Deny-list .....	52
Lab 1: Dynamic IP Deny-List .....	53
Labs - Dynamic Lazy Certificates .....	58
Lab 1: Use Dynamic Lazy Certificates from disk .....	59
Labs - Dynamic In-Memory Certificate .....	62
Lab 1: Use Dynamic In-memory Certificate.....	63
<b>Chapter 5 .....</b>	<b>71</b>
Labs - Configure NGINX to use HTTP/2.....	71
Lab 1: Configure NGINX to use HTTP/2 .....	72
Lab 2: Configure NGINX to not display version number on error pages .....	76
<b>Chapter 6 .....</b>	<b>79</b>
Labs - Secure API Traffic using JWT Authentication.....	79
Lab 1: Create auth_jwt Realm and Encode the Secret .....	80





# Chapter 1

---

## Labs - Introduction to Lab Environment

---

### Lab Contents

1. **Lab Environment**
2. **Login to the NGINX training system**
3. **Test Web Application**

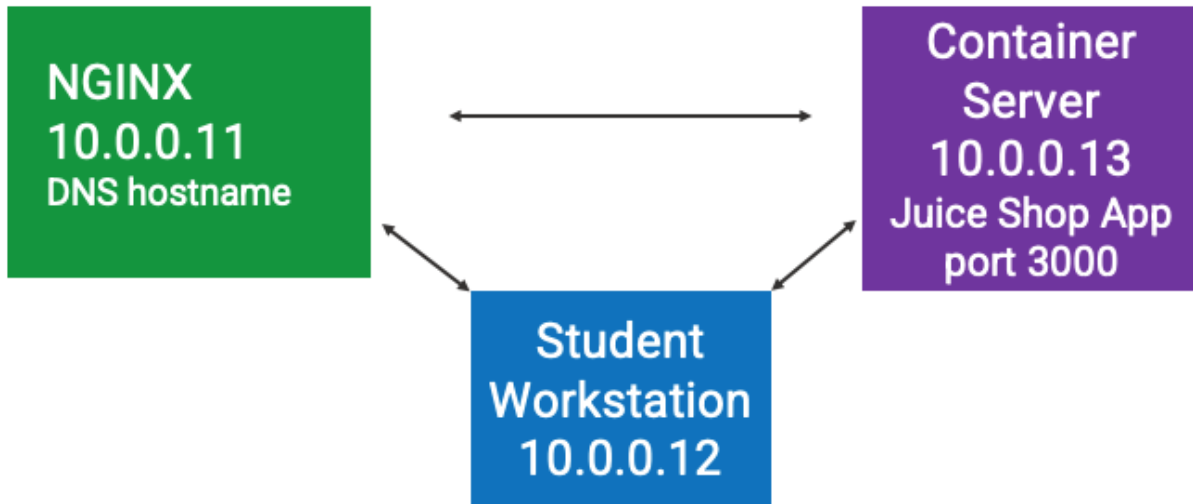
Estimated time to complete labs: [XX] minutes

### Lab Credentials and IP Address List

LAB SYSTEM	IP ADDRESS
NGINX	10.0.0.11
Student Workstation	10.0.0.12
Container Server	10.0.0.13

There are 3 lab systems: NGINX, Student Workstation and Container Server. All lab systems use the username **student** with the password **student**.

## Lab Network Topology



There are 3 lab systems: NGINX, Student Workstation and Container Server. The NGINX lab system has NGINX Plus installed. The Container Server system has the OWASP Juice Shop web application running on port 3000 using Kubernetes. The Student Workstation lab system is an Ubuntu Linux system that is a graphical system used to test access to the NGINX lab system and the Juice Shop application.



## Lab 1: Login to Your Training Machine

Estimated time for completion: **XX minutes**

### Requirements

The following tasks must be completed before beginning this lab:

- The OWASP Juice Shop web application must be running on the Server Container lab system and the NGINX lab system must have a `/etc/nginx/conf.d/juice.conf` file specifying `proxy_pass` and `upstream` to where the Juice Shop app is running.

### Overview

In this exercise, you practice logging into your NGINX lab system which is where most labs will be completed. Lab steps will explicitly say when to use the Student Workstation lab system instead of the NGINX lab system.

### Objectives

At the end of this lab you will be able to:

- Login to your NGINX lab system
- Verify the version of NGINX installed
- Test access to the Juice Shop web application directly on the Server Container lab system
- Test access to the Juice Shop web application via the NGINX lab system
- View the NGINX configuration file for the Juice Shop web application

### Steps

1. Log into your **NGINX** lab system. Use **student** as the *username* and as the *password*.
2. On the Desktop open the **Lab\_Files** folder and then open the **SNIPPETS.txt** file. This file lists the commands and NGINX code used in the lab steps and can be used to copy and paste lines into the command line interface (terminal) or browser. It's highly recommend that you copy/paste commands from the **SNIPPETS.txt** file.
3. All lab steps are done on the NGINX lab system unless explicitly stated.

4. Determine the version of NGINX.

```
$ nginx -v
```

5. Test the NGINX configuration files before you make any changes.

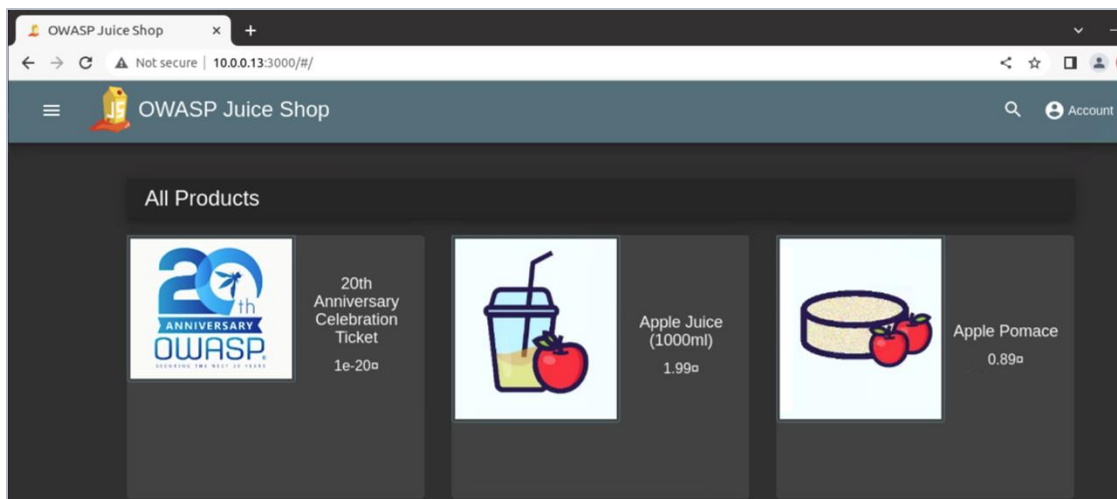
```
$ sudo nginx -t
```

6. View which NGINX processes are running. There should be one primary process and several worker processes.

```
$ ps aux | grep nginx
```

7. Open a Chrome browser to test accessing the Juice Shop web application directly by entering the IP address for your **Server Container** lab system (10.0.0.13) and port 3000 which is where the Juice Shop application is running.

```
http://10.0.0.13:3000
```



8. Open another tab in the Chrome browser and then enter the IP address for your **NGINX** lab system (10.0.0.11). You are accessing the Juice Shop application via NGINX rather than going directly to the Server Container lab system where the Juice Shop application is running.

```
http://10.0.0.11
```

9. Why do you not need to give port 3000 when accessing the Juice Shop application from the **NGINX** lab system?

10. To answer this, view the application configuration file **/etc/nginx/conf.d/juice.conf** looking at the **proxy\_pass** and **upstream** context where the upstream server is defined as **10.0.0.13:3000** which is the **Server Container** lab system where the Juice Shop application is running.

```
$ more /etc/nginx/conf.d/juice.conf
```

```
upstream juice_server {
    server 10.0.0.13:3000;
}

server {
    listen 80;
    root /usr/share/nginx/html;
    access_log /var/log/nginx/access.log combined;

    location / {
        proxy_pass http://juice_server;
    }

    error_page 500 502 503 504 /50x.html;

    location = /50x.html {
        root /usr/share/nginx/html;
    }
}
```

```
upstream juice_server {
    server 10.0.0.13:3000;
}

server {
    listen 80;
    root /usr/share/nginx/html;
    access_log /var/log/nginx/access.log combined;

    location / {
        proxy_pass http://juice_server;
    }

    error_page 500 502 503 504 /50x.html;

    location = /50x.html {
        root /usr/share/nginx/html;
    }
}
```

**NOTE**

From now on you will access the Juice Shop application using the NGINX lab system (10.0.0.11).

**Expected Results**

After logging into your NGINX lab system you will verify that NGINX Plus is already installed and that there is one NGINX master process along with one or more NGINX worker processes. You will test accessing the Juice Shop web application by going directly to the Server Container lab system using port 3000 and by using NGINX to access the Juice Shop application. You will view the contents of the `/etc/nginx/conf.d/juice.conf` file which has an upstream configured to access the Juice Shop application on the Server Container lab system.



## Chapter 2

---

### Labs - Setting up an HTTPS Server

---

#### Lab Contents

1. **Set DNS Variable**
2. **Generate Certs & Keys**
3. **Configure HTTPS Server**
4. **Configure SSL cipher, forward secrecy for Diffie-Hellman and test more secure NGINX configuration**

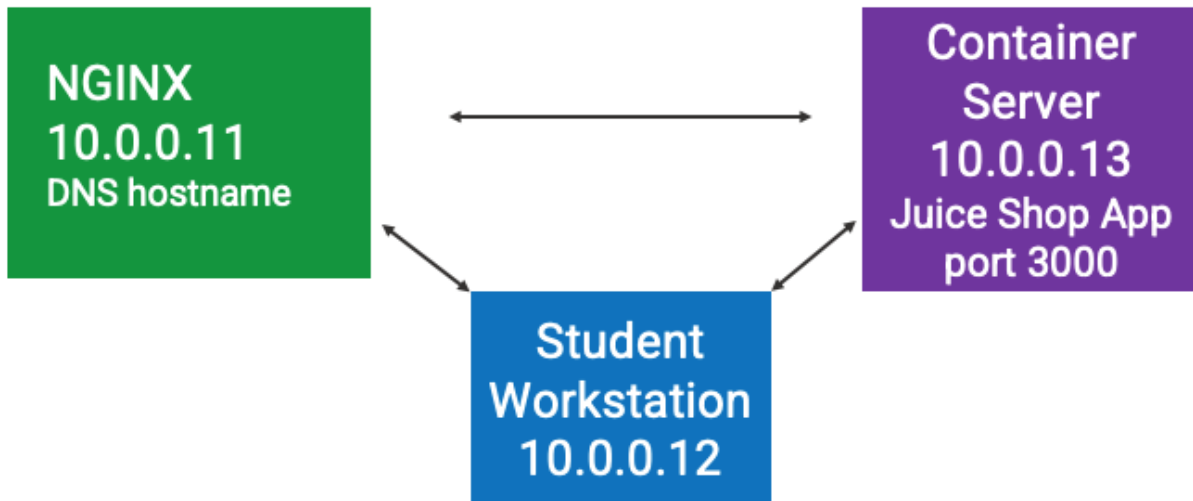
Estimated time to complete labs: [XX] minutes

#### Lab Credentials and IP Address List

LAB SYSTEM	IP ADDRESS
NGINX	10.0.0.11
Student Workstation	10.0.0.12
Container Server	10.0.0.13

## Lab Network Topology

### Lab Network Topology



There are 3 lab systems: NGINX, Student Workstation and Container Server. The NGINX lab system has NGINX Plus R28 installed. The Container Server system has the OWASP Juice Shop application running on port 3000 using Kubernetes. The Student Workstation lab system is an Ubuntu Linux system that is a graphical system used to test access to the NGINX lab system.

## Lab 1: Generating Certs and Keys

Estimated time for completion: **XX minutes**

### Requirements

### Overview

In this exercise, you generate certificates for the NGINX system in order to configure HTTPS for accessing the web application Juice Shop.

### Objectives

At the end of this lab you will be able to:

- Create a private CA Root Authority certificate
- Create a public and private certificate for the NGINX lab system
- Set permissions for cert files
- Add the NGINX directives for certificates

### Steps

1. Create a private key for the CA Root Authority, generate the X509 certificate for the CA Root Authority, and create the public and private certs for your NGINX system. The ***create\_certs*** script will create the CA Root Authority certs, and certs for a DNS name that represents your NGINX lab system ***www.nginxtraining.com***.

```
$ cd ~/ssl
$ less create_certs.sh
$ ./create_certs.sh www.nginxtraining.com
$ ls -ltr
```



#### NOTE

Using self-signed certificates or a certificate signed by a private CA is for training or testing purposes. In a production environment you would have certs signed by a legitimate CA Root Authority.

2. List the permissions, owner and group of the cert files created. Change the group to nginx for the newly created certificate files and change the permissions as follows.

```
$ sudo cp *cert www*key /etc/nginx/ssl
$ cd /etc/nginx/ssl
$ sudo chgrp nginx *
$ sudo chmod 640 *
$ ls -ltr
$ sudo file *
```

```
student@nginx:/etc/nginx/ssl$ ls -ltr
total 12
-rw-r----- 1 root nginx 1675 Jun 13 08:05 www.nginxtraining.com.key
-rw-r----- 1 root nginx 1310 Jun 13 08:05 www.nginxtraining.com.crt
-rw-r----- 1 root nginx 1375 Jun 13 08:05 ca-cert.crt
student@nginx:/etc/nginx/ssl$
student@nginx:/etc/nginx/ssl$ sudo file *
ca-cert.crt: PEM certificate
www.nginxtraining.com.crt: PEM certificate
www.nginxtraining.com.key: PEM RSA private key
student@nginx:/etc/nginx/ssl$
```



#### NOTE

The ca-cert.crt file is in PEM format. When using Dynamic certificate loading read access is required.

3. Before making changes, backup the original NGINX SSL configuration file /etc/nginx/ssl-configs/ssl-params.conf

```
$ sudo cp /etc/nginx/ssl-configs/ssl-params.{conf,orig}
```

4. Edit the ssl-params.conf file to add the **ssl\_certificate** directives.

```
$ sudo vim /etc/nginx/ssl-configs/ssl-params.conf
```

```
ssl_trusted_certificate /etc/nginx/ssl/ca-cert.crt;
ssl_certificate /etc/nginx/ssl/www.nginxtraining.com.crt;
ssl_certificate_key /etc/nginx/ssl/www.nginxtraining.com.key;
```



**NOTE**

The `ssl_trusted_certificate` directive needs to be before the `ssl_certificate` and `ssl_certificate_key` directives.

5. Save the file and exit the editor.
6. Test that the configuration syntax has no errors.

```
$ sudo nginx -t
```

**Expected Results**

You should now have the appropriate certificates set up for NGINX to access the Juice Shop application when HTTPS is configured next.



## Lab 2: Configuring an HTTPS Server

Estimated time for completion: **XX minutes**

### Requirements

The following tasks must be completed before beginning this lab:

- You must have a valid certificate and private key for your NGINX lab system.

### Overview

In this exercise, you configure HTTPS access to the Juice Shop app via NGINX

### Objectives

At the end of this lab you will be able to:

- Configure a server context listening on port 443 for HTTPS
- Set header values to be passed to the upstreams

### Steps

1. Verify **NGINX** has support for HTTPS.

```
$ sudo nginx -V
```

```
student@nginx:~$ sudo nginx -V
nginx version: nginx/1.23.2 (nginx-plus-r28)
built by gcc 9.3.0 (Ubuntu 9.3.0-10ubuntu2)
built with OpenSSL 1.1.1f  31 Mar 2020
TLS SNI support enabled
configure arguments: --prefix=/etc/nginx --sbin-path=/usr/sbin/nginx --modules-path=/usr/lib/nginx/modules --conf-path=/etc/nginx/nginx.conf --error-log-path=/var/log/nginx/error.log --http-log-path=/var/log/nginx/access.log --pid-path=/var/run/nginx.pid --lock-path=/var/run/nginx.lock --http-client-body-temp-path=/var/cache/nginx/client_temp --http-proxy-temp-path=/var/cache/nginx/proxy_temp --http-fastcgi-temp-path=/var/cache/nginx/fastcgi_temp --http-uwsgi-temp-path=/var/cache/nginx/uwsgi_temp --http-scgi-temp-path=/var/cache/nginx/scgi_temp --user=nginx --group=nginx --with-compat --with-file-aio --with-threads --with-http_addition_module --with-http_auth_request_module --with-http_dav_module --with-http_flv_module --with-http_gunzip_module --with-http_gzip_static_module --with-http_mp4_module --with-http_random_index_module --with-http_realip_module --with-http_secure_link_module --with-http_slice_module --with-http_ssl_module --with-http_stub_status_module --with-http_sub_module --with-http_v2_module --with-mail --with-mail_ssl_module --with-stream --with-stream_realip_module --with-stream_ssl_module --with-stream_ssl_preread_module --build=nginx-plus-r28 --with-http_auth_jwt_module --with-http_f4f_module --with-http_hls_module --with-http_session_log_module --with-http_proxy_protocol_vendor_module --with-stream_proxy_protocol_vendor_module --with-cc-opt='-g -O2 -fdebug-prefix-map=/data/builder/debuild/nginx-plus-1.23.2/debian/debuild-base/nginx-plus-1.23.2=. -fstack-protector-strong -Wformat -Werror=format-security -Wp,-D_FORTIFY_SOURCE=2 -fPIC' --with-ld-opt='-Wl,-Bsymbolic-functions -Wl,-z,relro -Wl,-z,now -Wl,--as-needed -pie'
student@nginx:~$
```

2. Before making any changes, backup the juice.conf configuration file.

```
$ sudo cp /etc/nginx/conf.d/juice.{conf,orig}
```

3. Edit the Juice Shop configuration file.

```
$ sudo vim /etc/nginx/conf.d/juice.conf
```

4. After the **listen** directive, add a **return** directive that forces all http traffic to **https** and close this server context by adding a right brace **}**. Do not delete any other lines in the file as we are not finished editing.

```
server {
    listen 80;
    return 301 https://$host$request_uri;
}
root /usr/share/nginx/html
```

5. After the close of the **return** directive, create a new server context listening on port **443**, with the **ssl** parameter enabled and make this the default server.

```
server {
    listen 80;
    return 301 https://$host$request_uri;
}
```

```
server {
    listen 443 ssl default_server;
    root /usr/share/nginx/html
```

6. After the **listen** directive, add the **server\_name** directive with the DNS name for your NGINX system, `www.nginxtraining.com`.

```
server {
    listen 443 ssl default_server;
    server_name www.nginxtraining.com;
    root /usr/share/nginx/html;
    access_log /var/log/nginx/access.log combined;
    ...
```

7. After the **root** directive, add an **include** directive for accessing your custom security settings file `/etc/nginx/ssl-configs/ssl-params.conf`.

```
server {
    listen 443 ssl default_server;
    server_name www.nginxtraining.com;
    root /usr/share/nginx/html;
    include /etc/nginx/ssl-configs/ssl-params.conf;
    access_log /var/log/nginx/access.log combined;
    ...
```

8. After the **access\_log** directive, add the following **proxy\_set\_header** directives.

```
server {
    listen 443 ssl default_server;
    server_name www.nginxtraining.com;
    root /usr/share/nginx/html;
    include /etc/nginx/ssl-configs/ssl-params.conf;
    access_log /var/log/nginx/access.log combined;
    proxy_set_header X-Real-IP $remote_addr;
    proxy_set_header X-Forwarded-For $proxy_add_x_forwarded_for;
    proxy_set_header Host $host;
```

9. Save the and exit the `juice.conf` configuration file.
10. Test your NGINX updated configuration files for any syntax errors.

```
$ sudo nginx -t
```

### **Expected Results**

You will now have a server context set up for HTTPS.



## Lab 3: Set SSL Ciphers

Estimated time for completion: **XX minutes**

### Requirements

The following tasks must be completed before beginning this lab:

- You must have a valid certificate and private key for your NGINX lab system.
- You must have a server block defined for Juice Shop to handle HTTPS requests.

### Overview

In this lab, edit the NGINX configuration file for the Juice Shop application to configure specific SSL ciphers.

### Objectives

At the end of this lab you will be able to:

- Specify which SSL cipher algorithms for NGINX to use.

### Steps

1. Use `openssl` to list the ciphers that match our request with a complete description of the protocol version, key exchange, authentication, encryption and mac algorithms used along with any key size restrictions.

```
$ openssl ciphers -v 'AES256+EECDH:AES256+EDH:!aNULL'
$ openssl ciphers -v | grep TLSv1.3
```

```

student@nginx:~$ openssl ciphers -v 'AES256+EECDH:AES256+EDH:!aNULL'
TLS_AES_256_GCM_SHA384 TLSv1.3 Kx=any Au=any Enc=AESGCM(256) Mac=AEAD
TLS_CHACHA20_POLY1305_SHA256 TLSv1.3 Kx=any Au=any Enc=CHACHA20/POLY1305(256) Mac=AEAD
TLS_AES_128_GCM_SHA256 TLSv1.3 Kx=any Au=any Enc=AESGCM(128) Mac=AEAD
ECDHE-ECDSA-AES256-GCM-SHA384 TLSv1.2 Kx=ECDH Au=ECDSA Enc=AESGCM(256) Mac=AEAD
ECDHE-RSA-AES256-GCM-SHA384 TLSv1.2 Kx=ECDH Au=RSA Enc=AESGCM(256) Mac=AEAD
ECDHE-ECDSA-AES256-CCM8 TLSv1.2 Kx=ECDH Au=ECDSA Enc=AESCCM8(256) Mac=AEAD
ECDHE-ECDSA-AES256-CCM TLSv1.2 Kx=ECDH Au=ECDSA Enc=AESCCM(256) Mac=AEAD
ECDHE-ECDSA-AES256-SHA384 TLSv1.2 Kx=ECDH Au=ECDSA Enc=AES(256) Mac=SHA384
ECDHE-RSA-AES256-SHA384 TLSv1.2 Kx=ECDH Au=RSA Enc=AES(256) Mac=SHA384
ECDHE-ECDSA-AES256-SHA TLSv1 Kx=ECDH Au=ECDSA Enc=AES(256) Mac=SHA1
ECDHE-RSA-AES256-SHA TLSv1 Kx=ECDH Au=RSA Enc=AES(256) Mac=SHA1
DHE-DSS-AES256-GCM-SHA384 TLSv1.2 Kx=DH Au=DSS Enc=AESGCM(256) Mac=AEAD
DHE-RSA-AES256-GCM-SHA384 TLSv1.2 Kx=DH Au=RSA Enc=AESGCM(256) Mac=AEAD
DHE-RSA-AES256-CCM8 TLSv1.2 Kx=DH Au=RSA Enc=AESCCM8(256) Mac=AEAD
DHE-RSA-AES256-CCM TLSv1.2 Kx=DH Au=RSA Enc=AESCCM(256) Mac=AEAD
DHE-RSA-AES256-SHA256 TLSv1.2 Kx=DH Au=RSA Enc=AES(256) Mac=SHA256
DHE-DSS-AES256-SHA256 TLSv1.2 Kx=DH Au=DSS Enc=AES(256) Mac=SHA256
DHE-RSA-AES256-SHA SSLv3 Kx=DH Au=RSA Enc=AES(256) Mac=SHA1
DHE-DSS-AES256-SHA SSLv3 Kx=DH Au=DSS Enc=AES(256) Mac=SHA1
student@nginx:~$ openssl ciphers -v | grep TLSv1.3
TLS_AES_256_GCM_SHA384 TLSv1.3 Kx=any Au=any Enc=AESGCM(256) Mac=AEAD
TLS_CHACHA20_POLY1305_SHA256 TLSv1.3 Kx=any Au=any Enc=CHACHA20/POLY1305(256) Mac=AEAD
TLS_AES_128_GCM_SHA256 TLSv1.3 Kx=any Au=any Enc=AESGCM(128) Mac=AEAD

```

**NOTE**

Remember to copy the command from the SNIPPETS.txt file or type in the quotes manually.



An "!" means that the cipher is deleted from the list. In this lab we are removing !aNULL which is the anonymous Diffie-Hellman algorithm which is vulnerable to MITM (Man In The Middle) attacks.

2. Edit the configuration file **ssl-params.conf** to add your allowed/preferred ciphers for both TLSv1.2 and TLSv1.3.

```
$ sudo vim /etc/nginx/ssl-configs/ssl-params.conf
```

```
ssl_ciphers AES256+EECDH:AES256+EDH:!aNULL;
```

```
ssl_conf_command Options PrioritizeChaCha;
```

```
ssl_conf_command Ciphersuites TLS_CHACHA20_POLY1305_SHA256;
```

```
ssl_protocols TLSv1.2 TLSv1.3;  
ssl_ciphers AES256+EECDH:AES256+EDH:!aNULL;  
ssl_conf_command Options PrioritizeChaCha;  
ssl_conf_command Ciphersuites TLS_CHACHA20_POLY1305_SHA256;  
ssl_prefer_server_ciphers on;
```

3. Save the and exit the juice.conf configuration file.
4. Test your NGINX updated configuration files for any syntax errors.

```
$ sudo nginx -t
```

### Expected Results

You should have 2 server blocks in your NGINX juice.conf file, one that listens on port 80 for HTTP traffic and one that listens on port 443 for HTTPS traffic. All HTTP traffic will be redirected to the server block for HTTPS.





## Lab 4: Configure Forward Secrecy

Estimated time for completion: **XX minutes**

### Requirements

The following tasks must be completed before beginning this lab:

- You must have a valid certificate and private key for your NGINX lab system.
- You must have a server block defined for Juice Shop to handle HTTPS requests.

### Objectives

At the end of this lab you will be able to:

- Configure Forward Secrecy in NGINX for Diffie-Hellman

### Steps

1. Generate a new Diffie-Hellman Key-Exchange. Use the `-dsaparam` flag in order to avoid testing for primality in all 4 million numbers.

```
$ sudo openssl dhparam -dsaparam -out /etc/nginx/dhparam.pem 4096
```

2. Edit the configuration file `ssl-params.conf` to add the `ssl_dhparam` directive with the `dhparam.pem` file you created in the previous step.

```
$ sudo vim /etc/nginx/ssl-configs/ssl-params.conf
```

```
ssl_dhparam /etc/nginx/dhparam.pem ;
```

3. Save the and exit the file.
4. Reload NGINX to use your updated configuration files.

```
$ sudo nginx -t && sudo nginx -s reload
```

5. Edit the file `/etc/hosts` to add the DNS name `www.nginxtraining.com` to be associated with your NGINX lab system `10.0.0.11`.

```
$ sudo vim /etc/hosts
```

```
10.0.0.11    www.nginxtraining.com
```

**NOTE**

Typically, if you are testing you would not edit system files but would instead use command options to handle your test case. For the curl command you can add the option **--resolve** **www.nginxtraining.com:443:10.0.0.11** but there may be limited information returned so for our training purposed adding the NGINX DNS name and IP address to /etc/hosts is preferable.

**Expected Results**

Your NGINX configuration now uses a specific SSL cipher algorithm and is set up to use Forward Secrecy with Diffie-Hellman.



## Lab 5: Testing HTTPS Configuration

Estimated time for completion: **XX minutes**

### Requirements

The following tasks must be completed before beginning this lab:

- You must have a valid certificate and private key for your NGINX lab system.
- You must have a server block defined for Juice Shop to handle HTTPS requests.

### Objectives

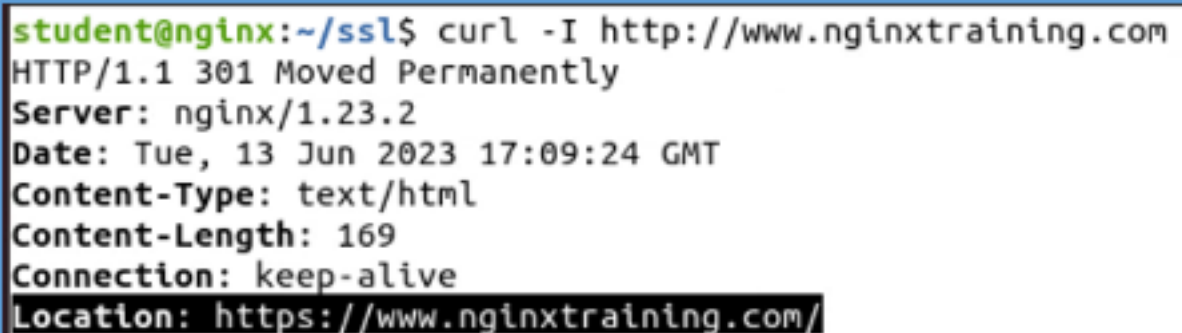
At the end of this lab you will be able to:

- Test HTTPS configuration

### Steps

1. Perform a curl test using HTTP. The output should have a 301 Moved Permanently message.

```
$ curl -I http://www.nginxtraining.com
```

A terminal window screenshot showing the output of the command 'curl -I http://www.nginxtraining.com'. The prompt is 'student@nginx:~/ssl\$'. The output is: 'HTTP/1.1 301 Moved Permanently', 'Server: nginx/1.23.2', 'Date: Tue, 13 Jun 2023 17:09:24 GMT', 'Content-Type: text/html', 'Content-Length: 169', 'Connection: keep-alive', and 'Location: https://www.nginxtraining.com/'.

```
student@nginx:~/ssl$ curl -I http://www.nginxtraining.com
HTTP/1.1 301 Moved Permanently
Server: nginx/1.23.2
Date: Tue, 13 Jun 2023 17:09:24 GMT
Content-Type: text/html
Content-Length: 169
Connection: keep-alive
Location: https://www.nginxtraining.com/
```

**NOTE**

This works for curl so you don't have to use the -k option, but in the browser it still gives ~~https~~ which means it's not secure/trusted. Depending on which browser you are using you can add your certificate to the browser trust but since in production you will not be using self-signed certs you can ignore these warnings in our training lab environment.

2. Perform another curl test using HTTP with the -L option to follow to the new location where the requested page has moved. These are two ways to not have curl complain about not being able to verify the certificate from using either a self-signed cert or like in our lab using a cert that was created from a private CA. The results of both commands are the same.

```
$ cd ~/ssl
$ curl -IL --cacert ca-cert.crt http://www.nginxtraining.com
$ curl -ILk http://www.nginxtraining.com
```

```
student@nginx:~/ssl$ curl -IL --cacert ca-cert.crt http://www.nginxtraining.com
HTTP/1.1 301 Moved Permanently
Server: nginx/1.23.2
Date: Tue, 13 Jun 2023 17:09:33 GMT
Content-Type: text/html
Content-Length: 169
Connection: keep-alive
Location: https://www.nginxtraining.com/

HTTP/1.1 200 OK
Server: nginx/1.23.2
Date: Tue, 13 Jun 2023 17:09:33 GMT
Content-Type: text/html; charset=UTF-8
Content-Length: 1924
Connection: keep-alive
Access-Control-Allow-Origin: *
X-Content-Type-Options: nosniff
X-Frame-Options: SAMEORIGIN
Feature-Policy: payment 'self'
Accept-Ranges: bytes
Cache-Control: public, max-age=0
Last-Modified: Tue, 13 Jun 2023 13:57:21 GMT
ETag: W/"784-188b50bfa5d"
Vary: Accept-Encoding
Strict-Transport-Security: max-age=63072000; includeSubdomains
X-Frame-Options: DENY
X-Content-Type-Options: nosniff
```

**NOTE**

Your output should have a 301 Moved Permanently but will also follow to the new location (200 OK) and that requires the certificate to be verified which is why we used the option `--cacert` to specify the CA Root Authority cert that our cert & key files were referenced with or the `-k` option to skip verification. Another way to avoid the verification issue is to put the CA cert you created into the system certificates store; for an Ubuntu system that would include copying the file `ca-cert.crt` to `/usr/local/share/ca-certificates` and then running `"sudo update-ca-certificates"`.

3. Perform a curl test using **HTTPS**. Your output should have a **200 OK** response.

```
$ curl -I https://www.nginxtraining.com
```

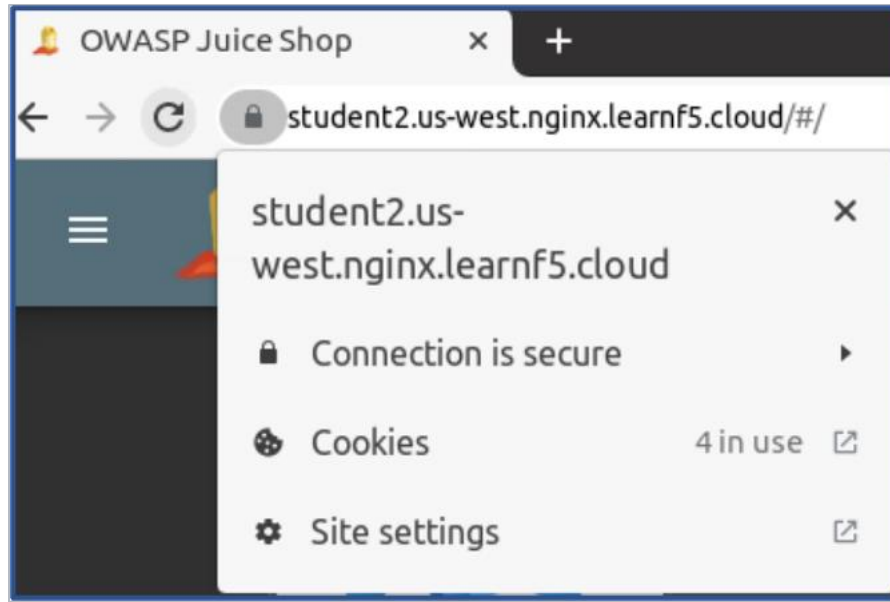
**TROUBLESHOOTING**

If you receive a 301 Moved Permanently response then make sure that the curl command is using **https** not **http**.

4. In your browser use **http://www.nginxtraining.com** to verify if the webpage now uses **https**. Because our certificate is signed by a private CA the browser can't verify which means the **https** will be crossed out and you need to click on the Advanced box at the bottom left and then click on the Proceed to `www.nginxtraining.com` (unsafe) link, which will take you to the Juice Shop web app via your NGINX lab system.

**NOTE**

If you were using a public (legitimate) CA certificate, then you would get a lock symbol or **https**.



### Expected Results

Your NGINX configuration now uses HTTPS with a specific SSL cipher algorithm and is set up to use Forward Secrecy with Diffie-Hellman.



---

## Labs - Secure Upstream Traffic

---

### Lab Contents

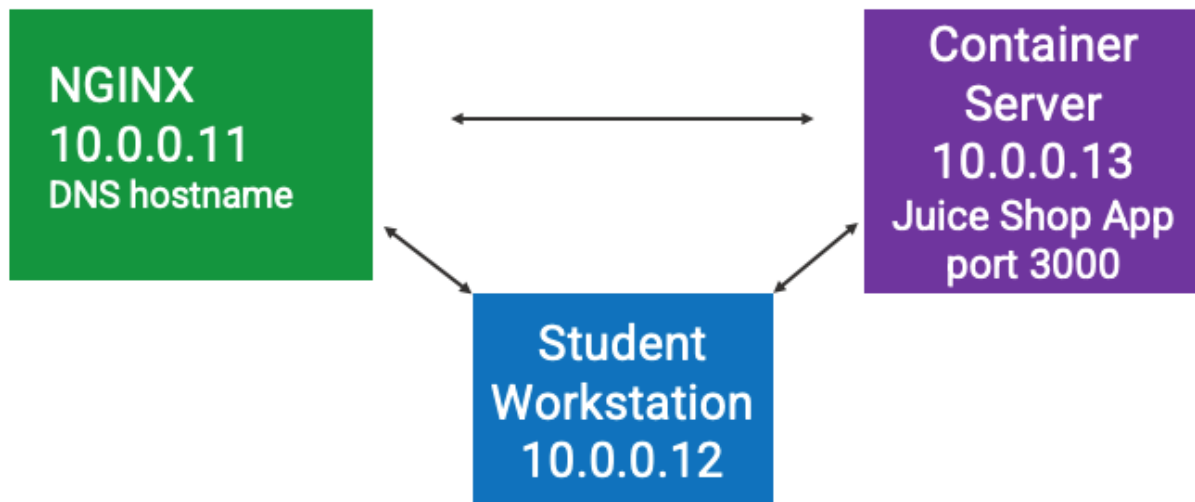
1. **Secure Upstream Traffic**
2. **Configure the NGINX Plus API Dashboard**
3. **Configure NGINX Access and Error Logging**

Estimated time to complete labs: [XX] minutes

### Lab Network Topology- Possibly Delete Topic?

#### Lab Network Topology

#### Lab Network Topology



There are 3 lab systems: NGINX, Student Workstation and Container Server. The NGINX lab system has NGINX Plus R28 installed. The Container Server system has the OWASP Juice Shop application running on port 3000 using Kubernetes. The Student Workstation lab system is an Ubuntu Linux system that is a graphical system used to test access to the NGINX lab system.

## Lab 1: Secure Upstream Traffic

Estimated time for completion: **XX minutes**

### Requirements

The following tasks must be completed before beginning this lab:

- You must have a valid certificate and private key for your NGINX lab system.
- You must have a server block defined for Juice Shop to handle HTTPS requests.

### Objectives

At the end of this lab you will be able to:

- Secure upstream traffic with ssl parameters
- Configure the NGINX API Dashboard to gather and view metrics
- Configure NGINX access and error logs

### Steps

1. Backup the configuration file **api\_server.conf**

```
$ sudo cp /etc/nginx/conf.d/api_server.{conf,orig}
```

2. Edit the configuration file **api\_server.conf**

```
$ sudo vim /etc/nginx/conf.d/api_server.conf
```

3. Enable HTTPS on port **8080** using the parameter **ssl**.

```
server {  
    listen 8080 ssl;  
    root /usr/share/nginx/html;  
    ...  
}
```

4. In the server block with ssl, create a **location** with the following proxy settings.

```
location / {  
    proxy_pass https://api_server;
```



```
proxy_ssl_trusted_certificate /etc/nginx/ssl/ca-cert.crt;
    proxy_ssl_verify off;
}
```

5. Save the `api_server.conf` file and then reload NGINX to use your updated configuration files.

```
$ sudo nginx -s reload
```

6. Test your configuration with curl. Your output should get a **502 Bad Gateway** response because you are trying to access upstreams that are not configured.

```
$ curl -k https://www.nginxtraining.com:8080/
```

7. View the contents of the **proxy-ssl-params.conf** file which defines the SSL protocols and HTTP version.

```
$ cat /etc/nginx/ssl-configs/proxy-ssl-params.conf
```

8. Edit the **api\_server.conf** file to add the **proxy-ssl-params.conf** file in the location / context.

```
$ sudo vim /etc/nginx/ssl-configs/proxy-ssl-params.conf
```

```
location / {
    include /etc/nginx/ssl-configs/proxy-ssl-params.conf;
    ...
}
```

9. Continue editing the `api_server.conf` file and for both server contexts, enable **ssl** and include the file **ssl-params.conf**

**NOTE**

This is **not the same pathname** as the previous step. There is no **proxy** in the name so be sure to edit the filename if you copy and paste the previous command or just copy from the Snippets file!

```
listen 8081 ssl;
include /etc/nginx/ssl-configs/ssl-params.conf;
...
```

```
listen 8082 ssl;  
include /etc/nginx/ssl-configs/ssl-params.conf;
```

10. Save the configuration file and reload NGINX.

```
$ sudo nginx -s reload
```

11. Test the upstream backend systems using the curl command.

```
$ cd ~/ssl  
$ curl --cacert ca-cert.crt https://www.nginxtraining.com:8080  
$ curl --cacert ca-cert.crt https://www.nginxtraining.com:8080  
$ curl --cacert ca-cert.crt https://www.nginxtraining.com:8080
```

#### NOTE



Your output should be like “**this is app1**”, then “**this is app2**” and then “**this is app1**” again. It does not matter which upstream server responds first since the default load balancing method is round robin so the backend servers should alternate serving the response. You might need to do more than three curl statements to see the Load Balancing start.

#### TROUBLESHOOTING:



If you are still getting a 502 Bad Gateway response message then make sure that you added "ssl" to both listen directives for ports 8081 and 8082 in the /etc/nginx/conf.d/api\_server.conf configuration file and that you reload to use the updated configuration file.

## Expected Results

The default round robin load balancing is set up for your application when using port 8080.



## Lab 2: Setup the NGINX Plus API Dashboard

Estimated time for completion: **XX minutes**

### Requirements

The following tasks must be completed before beginning this lab:

- NGINX load balancing using the default method, round robin, is configured when using port 8080.

### Overview

In this lab you will configure the NGINX Plus API Dashboard.

### Objectives

At the end of this lab you will be able to:

- Access the NGINX Plus API Dashboard.

### Steps

1. Edit the file `/etc/nginx/conf.d/api_server.conf`. In the context for 8080, create two new prefix locations: **api** and **dashboard**.

```
$ sudo vim /etc/nginx/conf.d/api_server.conf

location /api {
    api;
    access_log off;
}

location /dashboard {
    root /usr/share/nginx/html;
    try_files $uri $uri.html /dashboard.html;
}
```

**NOTE**

For our lab purposes we are turning off the access log since it generates a fair amount of information whereas in production you may not want to turn off.

2. Continue editing the **api\_server.conf** file, to gather status metrics add a **zone** directive in the upstream which is for the backends.

```
zone api_server_upstream 64k;
```

```
upstream api_server {  
    zone api_server_upstream 64k;  
  
    server 127.0.0.1:8081;  
    server 127.0.0.1:8082;  
}
```

3. Continue editing the **api\_server.conf** file, to gather status metrics by adding a **status\_zone** for the server listening on port 8080, which is for the front end, virtual servers. Save the file.

```
server {  
    listen 8080 ssl;  
    status_zone api_gateway;  
}
```

4. To gather more status metrics, edit the file `/etc/nginx/conf.d/juice.conf`. Add a **zone** directive in the upstream `juice_server`.

```
upstream juice_server {  
    server 10.0.0.13:3000;  
    zone juiceshop 64k;  
}
```

5. Continue editing the **juice.conf** file by adding a **status\_zone** for the server listening on port 443.

```
server {  
    listen 443 ssl default_server;  
    status_zone proxy;  
}
```

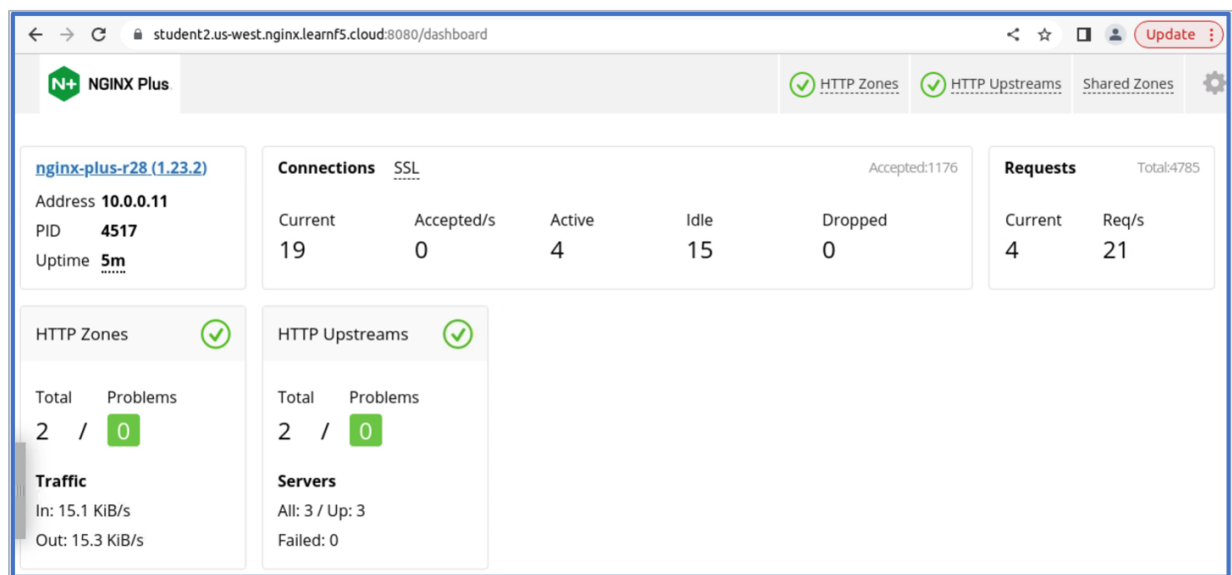
6. Save the file.
7. Test the syntax of your configuration files and when there are no errors reload the NGINX configuration.

```
$ sudo nginx -t && sudo nginx -s reload
```

- Send a few curl commands so that you have some newer data in your NGINX dashboard.

```
$ cd ~/ssl
$ curl --cacert ca-cert.crt https://www.nginxtraining.com:8080
$ curl --cacert ca-cert.crt https://www.nginxtraining.com:8080
```

- Enter <https://www.nginxtraining.com> in a browser. Note that you will still get a note from the browser about not trusted and this is okay since we have not added our private Root CA to the browsers trust.
- Access the NGINX dashboard from a browser using <https://www.nginxtraining.com:8080/dashboard>



#### NOTE



If the HTTP Upstreams is yellow, it should eventually turn to green. You could send more requests by doing several `curl -k https://www.nginxtraining.com:8080`

- In the NGINX dashboard click on the **HTTP Upstreams** tab to see both upstreams `api_server` and `juice_server`.
- In the dashboard click on the **HTTP Zones** tab to see the zones you added to the `api_server.conf` and `juice.conf` files (`api_gateway` and `proxy`).

13. In your dashboard click on the **Shared Zones** tab to see the `api_server_upstream` zone you added to the `api_server.conf` file and the `juiceshop` zone in the `juice.conf` file.
14. To get back to the original dashboard view, click the **NGINX Plus logo** (top left of the dashboard).

### Expected Results

The NGINX Plus Dashboard should be accessible to view upstreams, HTTP zones and shared zones.



## Lab 3: Configure NGINX Logging

Estimated time for completion: **XX minutes**

### Requirements

The following tasks must be completed before beginning this lab:

- You will need cert files and a server context for HTTPS
- You will need the `/etc/nginx/conf.d/juice.conf` file.

### Objectives

At the end of this lab you will be able to:

- Customize the format of the NGINX access log file.
- Access the NGINX access and error log files.

### Steps

1. Edit the configuration file `juice.conf`.

```
$ sudo vim /etc/nginx/conf.d/juice.conf
```

2. Add a custom log format in the http context (at the very top of the file, outside of any server context). Make sure to copy from the `SNIPPETS.txt` file.

```
log_format proxy_log "  
    Request: $request  
    Status: $status  
    Client: $remote_addr  
    Upstream: $upstream_addr  
    Forwarded-For: $proxy_add_x_forwarded_for  
    ssl_server_name: $ssl_server_name  
";  
  
upstream juice_server {  
    server 10.0.0.13:3000;  
    zone juiceshop 64k;  
}
```

3. Add the access and error log directives in the ssl server context and delete the original access\_log directive using the combined format. Then save the file.

```
access_log /var/log/nginx/juice.access.log proxy_log;  
error_log /var/log/nginx/juice.error.log info;
```

4. Test the syntax of your configuration files and when there are no errors reload the NGINX configuration.

```
$ sudo nginx -t && sudo nginx -s reload
```

5. Open another terminal to view the NGINX access log file.

```
$ sudo tail -f /var/log/nginx/juice.access.log
```

6. Test the new log format using curl commands and focus on the values output in the variable `ssl_server_name` which comes from the SNI TLS handshake. View the **curl\_script.sh** file to view the four curl commands and then run this script.

```
$ cat ~/curl_script.sh  
$ ~/curl_script.sh
```



```
Request: GET / HTTP/1.1
Status: 200
Client: 127.0.0.1
Upstream: 10.0.0.13:3000
Forwarded-For: 127.0.0.1
ssl_server_name: -

Request: GET / HTTP/1.1
Status: 200
Client: 127.0.0.1
Upstream: 10.0.0.13:3000
Forwarded-For: 127.0.0.1
ssl_server_name: localhost

Request: GET / HTTP/1.1
Status: 200
Client: 10.0.0.11
Upstream: 10.0.0.13:3000
Forwarded-For: 10.0.0.11
ssl_server_name: -

Request: GET / HTTP/1.1
Status: 200
Client: 10.0.0.11
Upstream: 10.0.0.13:3000
Forwarded-For: 10.0.0.11
ssl_server_name: www.nginxtraining.com
```

**NOTE**

When using these values the variable `ssl_server_name` is only set when passing `localhost` and `www.nginxtraining.com`.

7. Use Control-C to exit the tail command.

**Expected Results**

You will now have a custom log format for the NGINX access log.





## Chapter 3

---

### Labs - Set Limit Rates

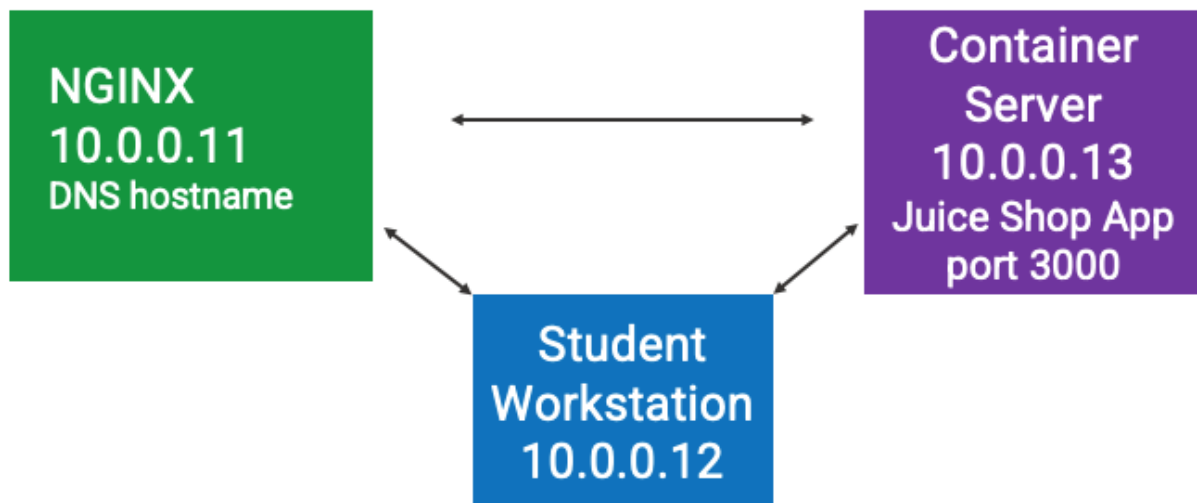
---

#### Lab Contents

1. **Configure geo limits**
2. **Configure Limit Rates**

Estimated time to complete labs: [XX] minutes

#### Lab Network Topology



## Lab 1: Configure NGINX limit rates

Estimated time for completion: **XX minutes**

### Requirements

The following tasks must be completed before beginning this lab:

- NGINX lab system with access to Juice Shop application running on another Linux lab system

### Objectives

At the end of this lab you will be able to:

- Use the NGINX geo and map directives.
- Configure limit request rates in NGINX.

### Steps

1. Create a copy of the juice.conf file before configuring limits in NGINX.

```
$ sudo cp /etc/nginx/conf.d/juice.{conf,PRElimits}
```

2. Edit the **juice.conf** file to add a **geo** and **map** directive in the http context (for example at the top of the file outside of any server context). The **map** directive sets the **\$limit\_key** variable to the client's IP address (in byte form via the variable **\$binary\_remote\_addr**).

```
$ sudo vim /etc/nginx/conf.d/juice.conf
```

```
geo $limit {  
    default 1;  
    127.0.0.1 0;  
}  
map $limit $limit_key {  
    0 "";  
    1 $binary_remote_addr;  
}
```

```

geo $limit {
    default 1;
    127.0.0.1 0;
}
map $limit $limit_key {
    0 "";
    1 $binary_remote_addr;
}

upstream juice_server {
    server 10.0.0.13:3000;
    zone juiceshop 64k;
}

```

3. In the http context, after the map, create a **limit\_req\_zone** using the **\$limit\_key** variable and set the request rate to 25 requests per second and then create a **limit\_conn\_zone** using the variable **\$server\_name** using 10 megabytes for the number of connections.

```

limit_req_zone $limit_key zone=req_zone:10m rate=25r/s;
limit_conn_zone $server_name zone=conn_zone:10m;

```

```

map $limit $limit_key {
    0 "";
    1 $binary_remote_addr;
}

limit_req_zone $limit_key zone=req_zone:10m rate=25r/s;
limit_conn_zone $server_name zone=conn_zone:10m;

```

4. In the ssl server context, set **limit\_req** with a burst of 15, nodelay and use **req\_zone** as the zone name. Also set the shared memory zone and the maximum allowed number of connections to 15 with the directive **limit\_conn**.

```

limit_req zone=req_zone burst=15 nodelay;
limit_conn conn_zone 15;

```

```
server {  
    listen 443 ssl default_server;  
    status_zone proxy;  
  
    limit_req zone=req_zone burst=15 nodelay;  
    limit_conn conn_zone 15;  
}
```

5. Save the juice.conf file and then reload your NGINX configuration to make it active.

```
$ sudo nginx -t && sudo nginx -s reload
```

**NOTE**

All of the changes you made to the juice.conf file are shown below (but this is not the whole file):

```
geo $limit {
    default 1;
    127.0.0.1 0;
}
map $limit $limit_key {
    0 "";
    1 $binary_remote_addr;
}

limit_req_zone $limit_key zone=req_zone:10m rate=25r/s;
limit_conn_zone $server_name zone=conn_zone:10m;

log_format proxy_log "
    Request: $request
    Status: $status
    Client: $remote_addr
    Upstream: $upstream_addr
    Forwarded-For: $proxy_add_x_forwarded_for
    ssl_server_name: $ssl_server_name
";

upstream juice_server {
    zone juiceshop 64k;
    server 10.0.0.13:3000;
}

server {
    listen 80;
    return 301 https://$host$request_uri;
}

server {
    listen 443 ssl default_server;
    status_zone proxy;
    limit_req zone=req_zone burst=15 nodelay;
    limit_conn conn_zone 15;
    server_name www.nginxtraining.com;
    include /etc/nginx/ssl-conflgs/ssl-params.conf;
    root /usr/share/nginx/html;
#    access_log /var/log/nginx/access.log combined;
    access_log /var/log/nginx/juice.access.log proxy_log;
    error_log /var/log/nginx/juice.error.log info;
```

**NOTE**

We will test these limits in the next lab exercise.

**Expected Results**

Your NGINX configuration has limit rates set for maximum requests.





---

## Labs - Logs

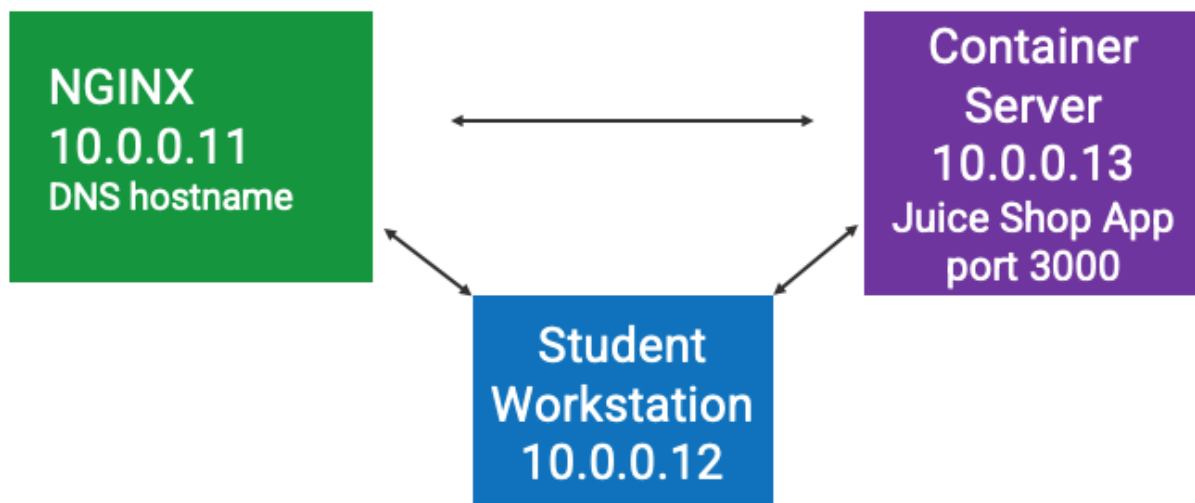
---

### Lab Contents

1. Access Log
2. Logging Limit Rates

Estimated time to complete labs: [XX] minutes

### Lab Network Topology



## Lab 1: NGINX Access and Error Log Files

Estimated time for completion: **XX minutes**

### Requirements

The following tasks must be completed before beginning this lab:

- Have limit request rates set in the juice.conf configuration file on the NGINX lab system.

### Objectives

At the end of this lab you will be able to:

- Test limits written to the NGINX access and error log files.

### Steps

6. In another terminal, view request information by using the tail -f command on the access log.

```
$ sudo tail -f /var/log/nginx/juice.access.log
```



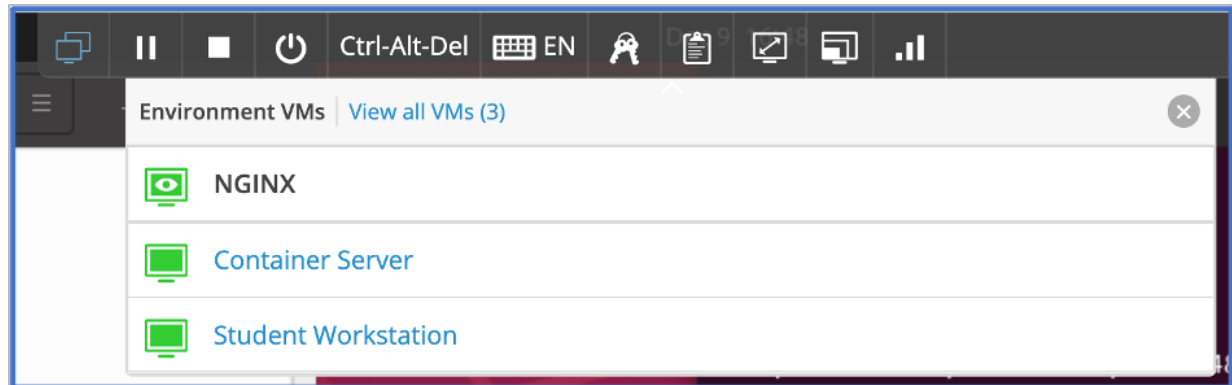
#### NOTE

When viewing log files, it is best to have a wide but not too tall terminal.

7. In another terminal, view the error information by using the tail -f command with the error log file.

```
$ sudo tail -f /var/log/nginx/juice.error.log
```

8. At the top of the Skytap environment, click on the two monitors so you can select the **Student Workstation** lab system.



9. Log in to the **Student Workstation** lab system, using **student** for the password.

10. On the Student Workstation lab system, open a terminal and run this **for** loop command by copying from the SNIPPETS.txt file.

```
$ for i in {0..50}; do (curl -kIs https://www.nginxtraining.com | head -n1 &) 2> /dev/null; done
```

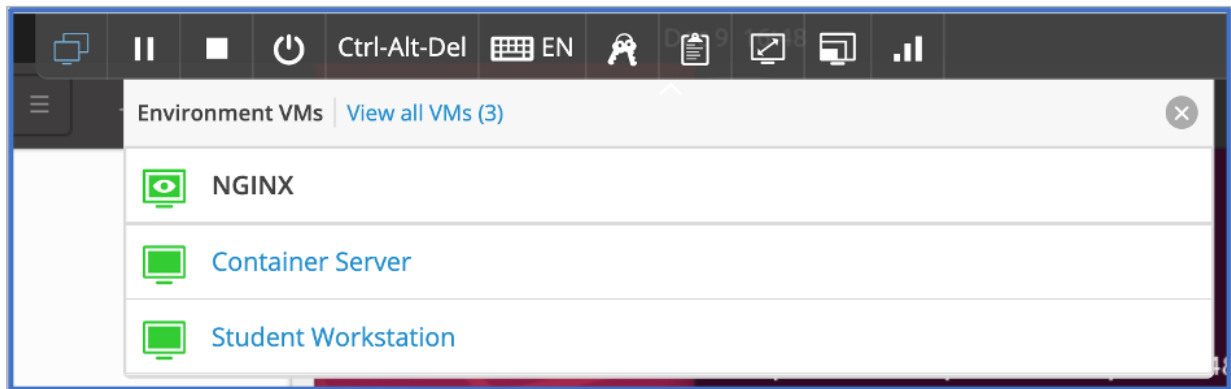
11. View the output from the for loop on your **Student Workstation** lab system:

```
student@jump:~$ for i in {0..50}; do (curl -kIs https://www.nginxtraining.com |  
head -n1 &) 2>/dev/null; done  
HTTP/1.1 200 OK  
HTTP/1.1 200 OK  
HTTP/1.1 200 OK  
HTTP/1.1 200 OK  
HTTP/1.1 200 OK  
HTTP/1.1 200 OK  
HTTP/1.1 200 OK  
HTTP/1.1 200 OK  
HTTP/1.1 200 OK  
HTTP/1.1 200 OK  
HTTP/1.1 200 OK  
HTTP/1.1 200 OK  
HTTP/1.1 200 OK  
HTTP/1.1 200 OK  
HTTP/1.1 200 OK  
HTTP/1.1 200 OK  
HTTP/1.1 200 OK  
HTTP/1.1 503 Service Temporarily Unavailable  
HTTP/1.1 200 OK  
HTTP/1.1 503 Service Temporarily Unavailable  
HTTP/1.1 503 Service Temporarily Unavailable  
HTTP/1.1 503 Service Temporarily Unavailable  
HTTP/1.1 503 Service Temporarily Unavailable  
HTTP/1.1 503 Service Temporarily Unavailable  
HTTP/1.1 200 OK
```

**NOTE**

The for loop sends more requests than what the limits are set to on your NGINX lab system (10.0.0.11) resulting in several ***Service Temporarily Unavailable*** messages.

12. At the top of the Skytap environment, click on the two monitors so you can select the **NGINX** lab system.



13. View the access log output from the **tail** command on your **NGINX** lab system.

```
Request: HEAD / HTTP/1.1
Status: 200
Client: 10.0.0.12
Upstream: 10.0.0.13:3000
Forwarded-For: 10.0.0.12
ssl_server_name: www.nginxtraining.com

Request: HEAD / HTTP/1.1
Status: 503
Client: 10.0.0.12
Upstream: -
Forwarded-For: 10.0.0.12
ssl_server_name: www.nginxtraining.com

Request: HEAD / HTTP/1.1
Status: 503
Client: 10.0.0.12
Upstream: -
Forwarded-For: 10.0.0.12
ssl_server_name: www.nginxtraining.com

Request: GET /socket.io/?EI0=4&transport=polling
Status: 200
Client: 10.0.0.11
Upstream: 10.0.0.13:3000
Forwarded-For: 10.0.0.11
ssl_server_name: www.nginxtraining.com
```

#### NOTE

You should get several **503** messages once the limit has been reached. The **Status: 200** indicates the request was successful. **Client: 10.0.0.12** indicates where the request originated, which is your Student Workstation lab system, where the for loop command was run. **Upstream: 10.0.0.13:3000** indicates where NGINX forwarded the request to the Server Container lab system where the Juice Shop application is running on port 3000.



14. Use the following command to search thru the juice.error.log file on your NGINX lab system to view several rate limit warning messages about excess requests.

```
$ sudo cat -vet /var/log/nginx/juice.error.log | grep limiting
```

```
2023/06/13 11:42:24 [error] 6950#6950: *2297 limiting requests, excess: 15.075 by zone "req_zone", client: 10.0.0.12,
server: www.nginxtraining.com, request: "HEAD / HTTP/1.1", host: "www.nginxtraining.com"$
2023/06/13 11:42:24 [error] 6949#6949: *2295 limiting requests, excess: 15.050 by zone "req_zone", client: 10.0.0.12,
server: www.nginxtraining.com, request: "HEAD / HTTP/1.1", host: "www.nginxtraining.com"$
2023/06/13 11:42:24 [error] 6948#6948: *2299 limiting requests, excess: 15.500 by zone "req_zone", client: 10.0.0.12,
server: www.nginxtraining.com, request: "HEAD / HTTP/1.1", host: "www.nginxtraining.com"$
```

#### NOTE

You should get several **503** messages once the limit has been reached. The **Status: 200** indicates the request was successful. **Client: 10.0.0.12** indicates where the request originated, which is your Student Workstation lab system, where the for loop command was run. **Upstream: 10.0.0.13:3000** indicates where NGINX forwarded the request to the Server Container lab system where the Juice Shop application is running on port 3000.



## Expected Results

After configuring rate limits in NGINX, you will have several errors and warning messages from sending more requests than allowed.



## Lab 2: Logging Limit Rates

Estimated time for completion: **XX minutes**

### Requirements

The following tasks must be completed before beginning this lab:

- NGINX configuration file set up with rate limits.

### Objectives

At the end of this lab you will be able to:

- Configure the NGINX limit request log level.
- Configure the NGINX limit request status response code.

### Steps

1. On the **NGINX** lab system, edit the juice.conf file.

```
$ sudo vim /etc/nginx/conf.d/juice.conf
```

2. In the ssl server context, add the **limit\_req\_log\_level** and the **limit\_req\_status**. Set the log\_level to warn when the error response code limit is exceeded.

```
limit_req_log_level warn;  
limit_req_status 4444;
```

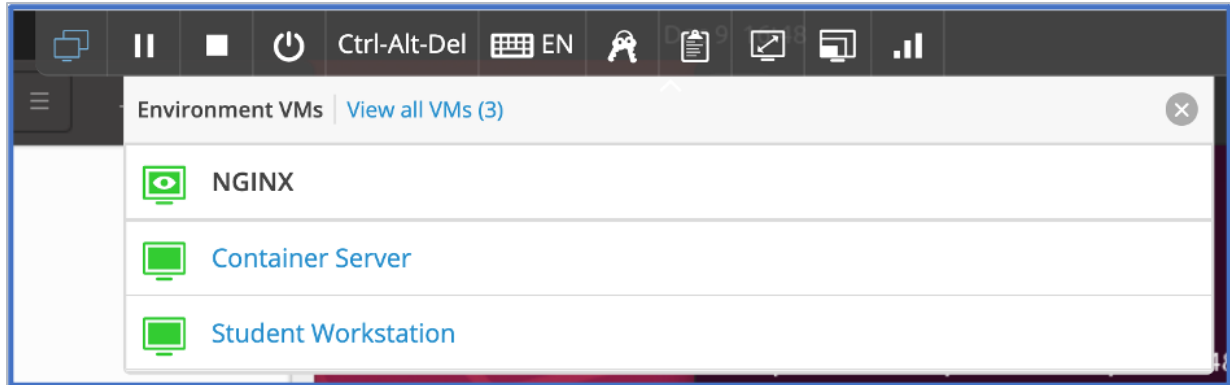
3. Save your configuration file and reload NGINX.

```
$ sudo nginx -s reload
```

4. View the access\_log file.

```
$ sudo tail -f /var/log/nginx/juice.access.log
```

5. At the top of the Skytap environment, click on the two monitors so you can select the **Student Workstation** lab system.



6. On the **Student Workstation** lab system, open a terminal and run this **for** loop command by copying from the SNIPPETS file.

```
$ for i in {0..50}; do (curl -kIs https://www.nginxtraining.com | head -n1 &) 2> /dev/null; done
```

7. View the output from the for loop on your Student Workstation lab system. There should only be HTTP/1.1 200 OK messages for the successful requests.
8. Switch back to your **NGINX** lab system to view the output from the NGINX access log. Note: You may have to scroll back to see the 444 status messages.

```
Request: HEAD / HTTP/1.1
Status: 200
Client: 10.0.0.12
Upstream: 10.0.0.13:3000
Forwarded-For: 10.0.0.12
ssl_server_name: -

Request: HEAD / HTTP/1.1
Status: 444
Client: 10.0.0.12
Upstream: -
Forwarded-For: 10.0.0.12
ssl_server_name: -

Request: HEAD / HTTP/1.1
Status: 444
Client: 10.0.0.12
Upstream: -
Forwarded-For: 10.0.0.12
ssl_server_name: -
```



**NOTE**

The output from the access log has the response code 444 which was set in our configuration file for when the error response code limit is exceeded. The 444 status closes the TCP connection to NGINX.

**TROUBLESHOOTING:**

If you do not get any **Status: 444** messages in the NGINX access log file, make sure to scroll back in the log file or use this command:

**grep 444 /var/log/nginx/juice.access.log**

9. Use control c to exit the tail command.
10. For class purposes go back to your juice.conf file that does not have limits by renaming juice.conf to juice.LIMITS and then rename the juice.PRElimits file to juice.conf.

```
$ sudo mv /etc/nginx/conf.d/juice.{conf,LIMITS}  
$ sudo mv /etc/nginx/conf.d/juice.{PRElimits,conf}
```

**Expected Results**

Configure and test logging limit rates by setting the limit request log level and setting the response status to 444 to terminate TCP connection.



---

## Labs - Setup Dynamic IP Deny-list

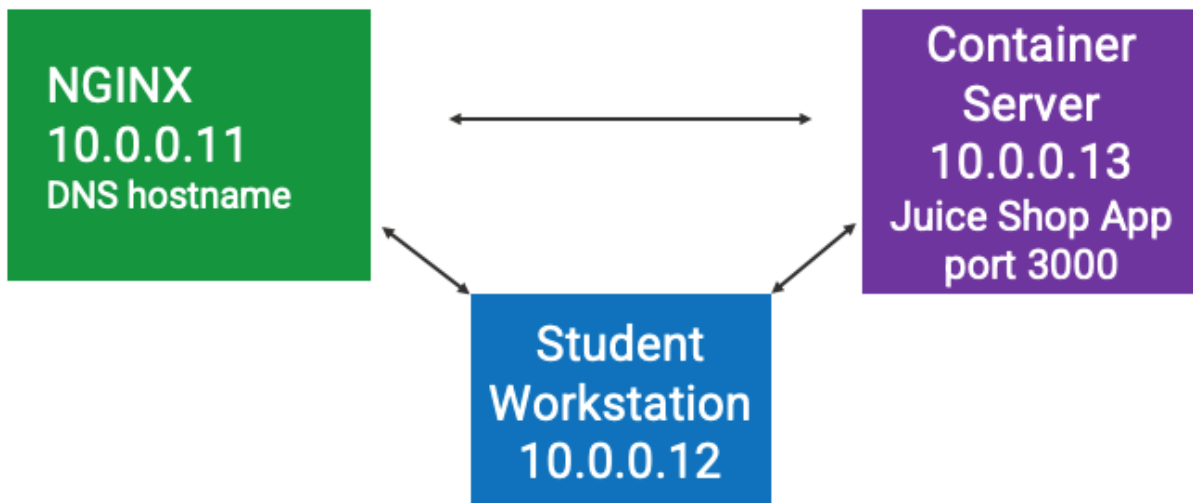
---

### Lab Contents

1. **Create Dynamic IP Deny-list**
2. **Set the \$target**
3. **Populate keyval Store**
4. **Test keyval\_zone**

Estimated time to complete labs: [XX] minutes

### Lab Network Topology



## Lab 1: Dynamic IP Deny-List

Estimated time for completion: **XX minutes**

### Requirements

The following tasks must be completed before beginning this lab:

- On the NGINX lab system, configure juice.conf file to access the Juice Shop application with HTTPS.

### Objectives

At the end of this lab you will be able to:

- Create a Dynamic IP Deny-list
- Use the NGINX Plus keyval store

### Steps

11. Create a directory that the nginx user has write access to store the key value file.

```
$ sudo mkdir /tmp/nginx  
$ sudo chown nginx:nginx /tmp/nginx
```

12. Edit the juice.conf configuration file.

```
$ sudo vim /etc/nginx/conf.d/juice.conf
```

13. In the http context add a **keyval\_zone** with the **state** parameter and add a keyval data store, mapping the client's ip address (\$remote\_addr) to the \$target variable.

```
keyval_zone zone=one:1m state=/tmp/nginx/one.keyval;  
keyval $remote_addr $target zone=one;
```

```
keyval_zone zone=one:1m state=/tmp/nginx/one.keyval;  
keyval $remote_addr $target zone=one;  
  
upstream juice_server {  
    server 10.0.0.13:3000;  
    zone juiceshop 64k;  
}
```

14. Continuing editing the juice.conf file, to create a new /api prefix in server 443 with write mode enabled on localhost only.

```
location /api {  
    api write=on;  
    allow 127.0.0.1;  
    deny all;  
}
```

15. Continuing editing the juice.conf file and after the api location add an **if** directive in the ssl server context with \$target as the argument and a return status 403.

```
if ($target) {  
    return 403;  
}
```



#### NOTE

All of the changes you made to the juice.conf file are shown below (but this is not the whole file):

```
proxy_set_header X-Real-IP $remote_addr;  
proxy_set_header X-Forwarded-For $proxy_add_x_forwarded_for;  
proxy_set_header Host $host;  
  
location /api {  
    api write=on;  
    allow 127.0.0.1;  
    allow 10.0.0.11;  
    deny all;  
}  
  
if ($target) {  
    return 403;  
}  
  
location / {  
    proxy_pass http://juice_server;  
}
```

16. Save your configuration file and reload NGINX.

```
$ sudo nginx -t && sudo nginx -s reload
```

17. Send a POST request to populate the keyval database with your **Student Workstation** lab systems IP address (10.0.0.12). By setting a value of “1” the Student Workstation lab system will be on the deny-list. You should get a **201 Created** response code which means your curl command was successful. Remember to copy the curl command from the SNIPPETS text file.

```
$ curl -k -v -H "Content-Type: application/json" -d '{"10.0.0.12": "1"}'
'https://127.0.0.1/api/8/http/keyvals/one'
```

#### NOTE

You should get a “201 Created” response code, which means your curl command worked.



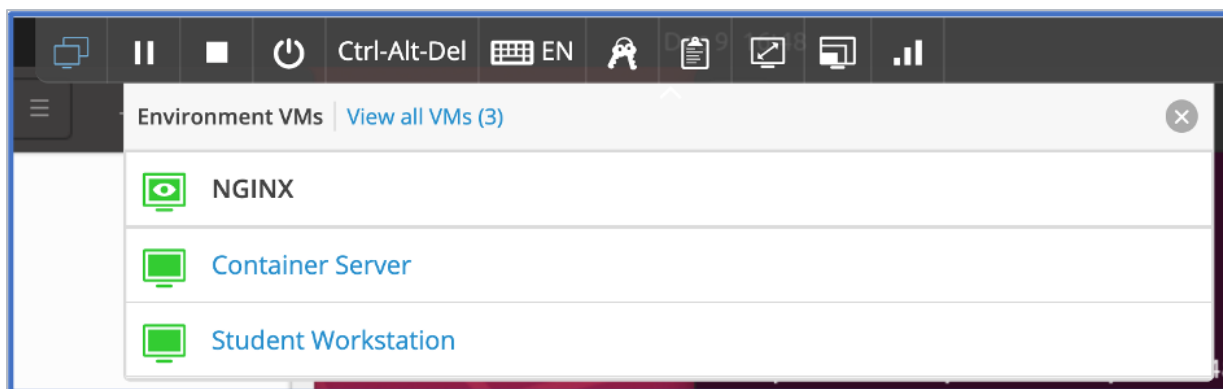
#### TROUBLESHOOTING:

If you get a 409 Conflict message showing the KeyvalKeyExists then you already have a value in memory (test with the next lab step) and you can continue or you could clear out the data and do this command again.

18. Send a GET request to confirm it works. Your output should be: {"10.0.0.12": "1"}

```
$ curl -k https://127.0.0.1/api/8/http/keyvals/one
```

19. At the top of the Skytap environment, click on the two monitors so you can select the **Student Workstation** lab system.



20. On the **Student Workstation** lab system, open a browser and enter your **NGINX** lab systems IP address 10.0.0.11. You should get a “**403 Forbidden**” response code

since the data part of the key-value pair is set to “1” which means “true” so NGINX will deny your Student Workstation lab system since it is on the deny-list.

<https://www.nginxtraining.com>



#### TROUBLESHOOTING:

You may have to click on some browser security questions to get to the 403 Forbidden response code: clear your browser data cache or use an incognito or private browser window.

21. Switch back to the **NGINX** lab system to use the NGINX api to update the **keyval\_zone** data to change the value to a zero. You should get a "204 No Content" response.

```
$ curl -k -i -X PATCH -d '{"10.0.0.12": "0"}' -s 'https://127.0.0.1/api/8/http/keyvals/one'
```

```
student@nap:~$ curl -k -i -X PATCH -d '{"10.0.0.12": "0"}' -s 'https://127.0.0.1/api/4/http/keyvals/one'
HTTP/1.1 204 No Content
Server: nginx
Date: Wed, 21 Dec 2022 13:55:33 GMT
Connection: keep-alive
Strict-Transport-Security: max-age=63072000; includeSubdomains
X-Frame-Options: DENY
X-Content-Type-Options: nosniff
```

22. On the NGINX lab system, send a GET request to confirm it updated the value from “1” to “0”.

```
$ curl -k https://127.0.0.1/api/8/http/keyvals/one
```

Your output should be:

```
{"10.0.0.12": "0"}
```

23. Switch back to the **Student Workstation** lab system and in the browser, refresh the webpage <https://10.0.0.11> to view if the result is different than previously?

**NOTE**

Yes. You should now be able to get to the Juice Shop web page without a Forbidden error message because the “0” means false, i.e., don’t denylist things from this IP address.

24. Switch back to the NGINX lab system and use the NGINX api to update the **keyval\_zone** data to **DELETE** the entry. You should get a 204 No Content response.

```
$ curl -k -i -X DELETE -d '{"10.0.0.12": "0"}' -s  
'https://127.0.0.1/api/8/http/keyvals/one'
```

25. Verify that the value has been deleted.

```
$ curl -k https://127.0.0.1/api/4/http/keyvals/one
```

Your output should be:

```
{}
```

**Expected Results**

You have tested using a deny-list in your NGINX configuration.



---

## Labs - Dynamic Lazy Certificates

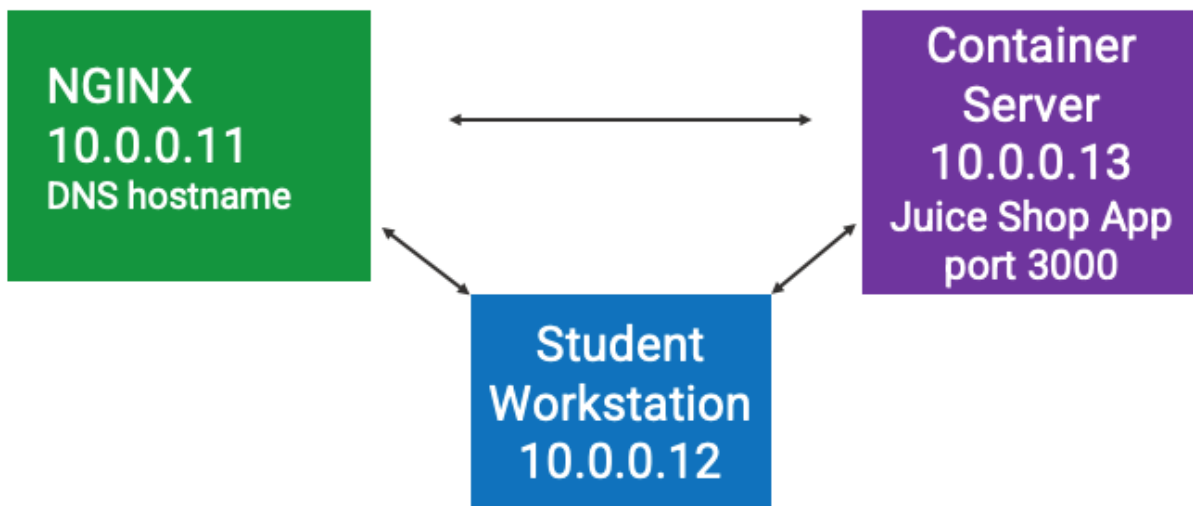
---

### Lab Contents

1. Use Dynamic Lazy Certificate Loading from disk

Estimated time to complete labs: [XX] minutes

### Lab Network Topology





## Lab 1: Use Dynamic Lazy Certificates from disk

Estimated time for completion: **XX minutes**

### Requirements

The following tasks must be completed before beginning this lab:

- On your NGINX lab system a juice.conf file with an upstream to the Juice Shop application with certs and HTTPS set up.

### Objectives

At the end of this lab you will be able to:

- Use NGINX Dynamic Lazy Certificate file loading from disk.

### Steps

1. Before changing to use the NGINX Dynamic Certificate loading method, backup the configuration files: ssl-params.conf, juice.conf, and api\_server.conf

```
$ sudo cp /etc/nginx/ssl-configs/ssl-params.{conf,traditional}  
$ sudo cp /etc/nginx/conf.d/juice.{conf,traditional}  
$ sudo cp /etc/nginx/conf.d/api_server.{conf,traditional}
```

2. Edit the **ssl-params.conf** file to add new ssl certificate directives to use Dynamic Lazy loading. These directives use special variables for dynamically loading the appropriate certificate based on SNI.

```
$ sudo vim /etc/nginx/ssl-configs/ssl-params.conf
```

```
ssl_certificate /etc/nginx/ssl-configs/certs/$ssl_server_name.crt;  
ssl_certificate_key /etc/nginx/ssl-configs/certs/$ssl_server_name.key;
```



#### NOTE

The ssl\_trusted\_certificate can't be loaded dynamically.

3. Comment out the original 2 certificate directives that use the traditional method of listing each certificate file and pathname on disk. The beginning of the ssl-params.conf file is shown.

```
ssl_trusted_certificate /etc/nginx/ssl/ca-cert.crt;

#Dynamic Lazy Loading
ssl_certificate /etc/nginx/ssl/$ssl_server_name.crt;
ssl_certificate_key /etc/nginx/ssl/$ssl_server_name.key;

#Traditional cert loading
ssl_certificate /etc/nginx/ssl/www.nginxtraining.com.crt;
ssl_certificate_key /etc/nginx/ssl/www.nginxtraining.com.key;
```

4. Test and reload NGINX to use your updated configuration files.

```
$ sudo nginx -s reload
```

5. Test your configuration with curl. Make sure to scroll back to the beginning output to see the connection, port information and response 200 OK.

```
$ cd ~/ssl
```

```
$ curl -sIvk --cacert ca-cert.crt https://www.nginxtraining.com
```

```

student@nginx:~/ssl$ curl -sIv --cacert ca-cert.crt https://www.nginxtraining.com
* Trying 10.0.0.11:443...
* TCP_NODELAY set
* Connected to www.nginxtraining.com (10.0.0.11) port 443 (#0)
* ALPN, offering h2
* ALPN, offering http/1.1
* successfully set certificate verify locations:
*   CAfile: ca-cert.crt
*   CApath: /etc/ssl/certs
* TLSv1.3 (OUT), TLS handshake, Client hello (1):
* TLSv1.3 (IN), TLS handshake, Server hello (2):
* TLSv1.3 (OUT), TLS change cipher, Change cipher spec (1):
* TLSv1.3 (OUT), TLS handshake, Client hello (1):
* TLSv1.3 (IN), TLS handshake, Server hello (2):
* TLSv1.3 (IN), TLS handshake, Encrypted Extensions (8):
* TLSv1.3 (IN), TLS handshake, Certificate (11):
* TLSv1.3 (IN), TLS handshake, CERT verify (15):
* TLSv1.3 (IN), TLS handshake, Finished (20):
* TLSv1.3 (OUT), TLS handshake, Finished (20):
* SSL connection using TLSv1.3 / TLS_CHACHA20_POLY1305_SHA256
* ALPN, server accepted to use http/1.1
* Server certificate:
*   subject: CN=www.nginxtraining.com; OU=Training; O=F5; L=SanFrancisco; ST=CA; C=US
*   start date: Jun 13 14:57:08 2023 GMT
*   expire date: May 20 14:57:08 2123 GMT
*   subjectAltName: host "www.nginxtraining.com" matched cert's "www.nginxtraining.com"
*   issuer: CN=www.nginxtraining.com-ca; OU=Training; O=F5; L=SanFrancisco; ST=CA; C=US
*   SSL certificate verify ok.
> HEAD / HTTP/1.1
> Host: www.nginxtraining.com
> User-Agent: curl/7.68.0
> Accept: */*
>
* TLSv1.3 (IN), TLS handshake, Newsession Ticket (4):
* TLSv1.3 (IN), TLS handshake, Newsession Ticket (4):
* old SSL session ID is stale, removing
* Mark bundle as not supporting multiuse
< HTTP/1.1 200 OK
HTTP/1.1 200 OK
< Server: nginx/1.23.2

```

## Expected Results

You have configured NGINX to use dynamic lazy certificate loading instead of the traditional certificate loading method.



---

## Labs - Dynamic In-Memory Certificate

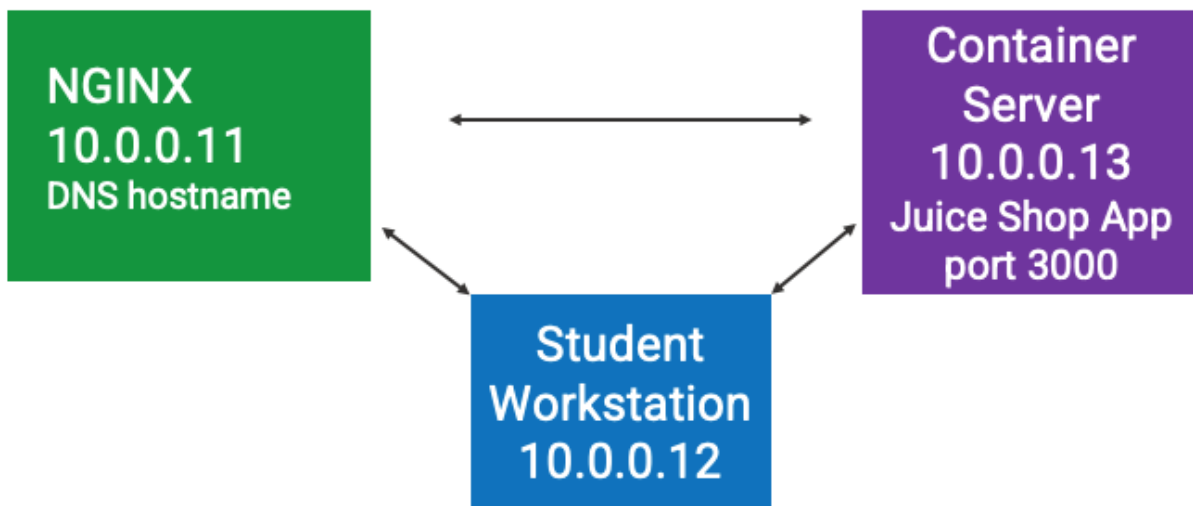
---

### Lab Contents

#### 1. Using NGINX Dynamic In-memory Certificates

Estimated time to complete labs: [XX] minutes

### Lab Network Topology



## Lab 1: Use Dynamic In-memory Certificate

Estimated time for completion: **XX minutes**

### Requirements

The following tasks must be completed before beginning this lab:

- On your NGINX lab system a juice.conf file with an upstream to the Juice Shop application.

### Objectives

At the end of this lab you will be able to:

- Configure and test using NGINX Dynamic In-memory certificates.

### Steps

1. Before changing to use the NGINX Dynamic In-memory certificate loading method, backup the following configuration files: **ssl-params.conf**, **api\_server.conf**, and **juice.conf**.

```
$ sudo cp /etc/nginx/ssl-configs/ssl-params.{conf,lazy}
$ sudo cp /etc/nginx/conf.d/api_server.{conf,lazy}
$ sudo cp /etc/nginx/conf.d/juice.{conf,lazy}
```

2. Edit the **juice.conf** file in the http context (top of file) add two new **keyval\_zone** directives; one for the certificate data and one for the private key data and then add two new **keyval** directives mapping the key values and zones. Save the juice.conf file.

```
keyval_zone zone=ssl_cert:1m;
keyval_zone zone=ssl_key:1m;
keyval $ssl_server_name $crt_pem zone=ssl_cert;
keyval $ssl_server_name $key_pem zone=ssl_key;
```

```

keyval_zone zone=ssl_crt:1m;
keyval_zone zone=ssl_key:1m;
keyval $ssl_server_name $crt_pem zone=ssl_crt;
keyval $ssl_server_name $key_pem zone=ssl_key;

log_format proxy_log "
    Request: $request
    Status: $status
    Client: $remote_addr
    Upstream: $upstream_addr
    Forwarded-For: $proxy_add_x_forwarded_for
    ssl_server_name: $ssl_server_name
";

keyval_zone zone=one:1m state=/tmp/nginx/one.keyval;
keyval $remote_addr $target zone=one;

upstream juice_server {
    zone juiceshop 64k;
    server 10.0.0.13:3000;
}

```

3. Reload NGINX to use your updated configuration files

```
$ sudo nginx -t && sudo nginx -s reload
```

4. Configure the certificate and key files into a format that NGINX can store in memory and that your student account has access to. View the **upload.sh** script file.

```
$ cd ~/ssl
$ cat upload.sh
```

```

student@nginx:~/ssl$ cat upload.sh
#!/bin/bash

HOST=$1
echo HOST: $HOST
for EXT in key crt ; do
    CERT=$(sed 's/$/\n/' $HOST.$EXT | tr -d '\n')
    URL=https://$HOST/apl/8/http/keyvals/ssl_$EXT
    echo {"$HOST": {"$CERT"}} | curl -kvH "Content-Type: application/json" -d @- $URL
done

```

5. Run the script **upload.sh** passing it the DNS name of your NGINX system. You must run this script while in the ~/ssl directory so it has access to the cert files.

```
$ ./upload.sh www.nginxtraining.com
```

There should be 2 of these “201 Created” lines - one for the key file and one for the cert file.

If you look closely at the output there is a **POST** HTTP request which shows where the value is being written.

```
POST /api/8/http/keyvals/ssl_cert HTTP/1.1
```

Also, the **Location** line contains the URL to the written memory file which you can retrieve using **curl -k Location\_URL\_value**.

```
< HTTP/1.1 201 Created
< Server: nginx/1.23.2
< Date: Tue, 13 Jun 2023 23:36:27 GMT
< Content-Length: 0
< Location: https://www.nginxtraining.com/api/8/http/keyvals/ssl_cert/
< Connection: keep-alive
< Strict-Transport-Security: max-age=63072000; includeSubdomains
< X-Frame-Options: DENY
< X-Content-Type-Options: nosniff
```

### Troubleshooting:

Make sure you run the UPLOAD script in the terminal where you have the nginx group as part of your student account (where you ran `exec su -l student`) and make sure you are in the directory `/etc/nginx/ssl-configs/certs`.

If you need to delete the key value store memory entry and run the UPLOAD script again do these commands:

```
$ curl -k -i -X DELETE https://$DNS/api/6/http/keyvals/ssl_key
$ curl -k -i -X DELETE https://$DNS/api/6/http/keyvals/ssl_cert
```

OR

```
$ sudo systemctl stop nginx
$ sudo systemctl status nginx
$ sudo systemctl start nginx
```

6. To test that both the certificate and private key files are in memory in the keyval store location.







**NOTE**

Your file should look like:



```
ssl_trusted_certificate /etc/nginx/ssl/ca-cert.crt;

#Dynamic InMemory Cert Loading
ssl_certificate data:$crt_pem;
ssl_certificate_key data:$key_pem;

#Dynamic Lazy Loading
#ssl_certificate /etc/nginx/ssl/$ssl_server_name.crt;
#ssl_certificate_key /etc/nginx/ssl/$ssl_server_name.key;

#Traditional cert loading
ssl_certificate /etc/nginx/ssl/www.nginxtraining.com.crt;
ssl_certificate_key /etc/nginx/ssl/www.nginxtraining.com.key;
```

NOTE: Leave the `ssl_trusted_certificate` directive as it is since it is not loaded dynamically.

8. Edit the `api_server.conf` file to add proxy ssl name information using the directives **`proxy_ssl_name`** and **`proxy_ssl_server_name`**. Add these to the 8080 server block under the location `/` and after your `proxy_pass` directive.

```
proxy_ssl_name www.nginxtraining.com;
proxy_ssl_server_name on;
```

```
$ sudo vim /etc/nginx/conf.d/api_server.conf
```

```
location / {
    include /etc/nginx/ssl-configs/proxy-ssl-params.conf;
    proxy_pass https://api server;
    proxy_ssl_name student6.us-west.nginx.learnf5.cloud;
    proxy_ssl_server_name on;
    proxy_ssl_trusted_certificate /etc/letsencrypt/live/student6.us-west.nginx.l
earnf5.cloud/fullchain.pem;
    proxy_ssl_verify off;
}
```

9. Test and reload NGINX to use your updated configuration files.

```
$ sudo nginx -t && sudo nginx -s reload
```

10. Send a GET request to confirm Dynamic Certificate loading is working. The output will be your private key and cert.

```
$ curl -k https://www.nginxtraining.com:8080
$ curl -k https://www.nginxtraining.com:8080
$ curl -k https://www.nginxtraining.com:8080
```

**NOTE**

You should get alternating upstreams responding such as “this is app1”, then “this is app2!” followed by “this is app1”.

11. Test access to the Juice Shop app.

```
$ curl -k https://www.nginxtraining.com
```

12. For our lab purposes revert back to using the traditional certificates by copying your backup files back to the original files. This will now use the traditional certificate.

```
$ sudo cp /etc/nginx/ssl-configs/ssl-params.{conf,dynamic-in-memory}
$ sudo cp /etc/nginx/ssl-configs/ssl-params.{traditional,conf}
$ sudo cp /etc/nginx/conf.d/juice.{conf,dynamic-in-memory}
$ sudo cp /etc/nginx/conf.d/api_server.{traditional,conf}
$ sudo cp /etc/nginx/conf.d/api_server.{conf,dynamic-in-memory} $ sudo cp
/etc/nginx/conf.d/juice.{traditional,conf}
```

13. Test and reload NGINX to use your updated configuration files.

```
$ sudo nginx -t && sudo nginx -s reload
```

14. Test your configuration with curl. Make sure to scroll back to the beginning output to see 200 OK message.

```
$ cd ~/ssl
$ curl -LIv --cacert ca-cert.crt https://www.nginxtraining.com
```

15. Test using a browser. Your output should be either “this is app1” or “this is app2” and then the Juice Shop app.

```
https://www.nginxtraining.com:8080
https://www.nginxtraining.com
```

**Expected Results**

You configure and test using the NGINX Dynamic In-memory certificates.





## Chapter 5

---

### Labs - Configure NGINX to use HTTP/2

---

#### Lab Contents

1. **Configure NGINX to use HTTP/2**
2. **Configure NGINX to not display the version number on error pages**

Estimated time to complete labs: [XX] minutes

## Lab 1: Configure NGINX to use HTTP/2

Estimated time for completion: **XX minutes**

### Requirements

The following tasks must be completed before beginning this lab:

- Have the configuration file `/etc/nginx/conf.d/api_server.conf`

### Objectives

At the end of this lab you will be able to:

- Configure NGINX to use HTTP/2.

### Steps

16. Edit the **api\_server.conf** file.

```
$ sudo vim /etc/nginx/conf.d/api_server.conf
```

17. In the server context for 8080, add **http2** to the listen directive. Save the file.

```
server {  
    listen 8080 ssl http2;  
    ...  
}
```

18. Reload NGINX.

```
$ sudo nginx -s reload
```

19. Test using a browser. Your output should be either “this is app1” or “this is app2”.

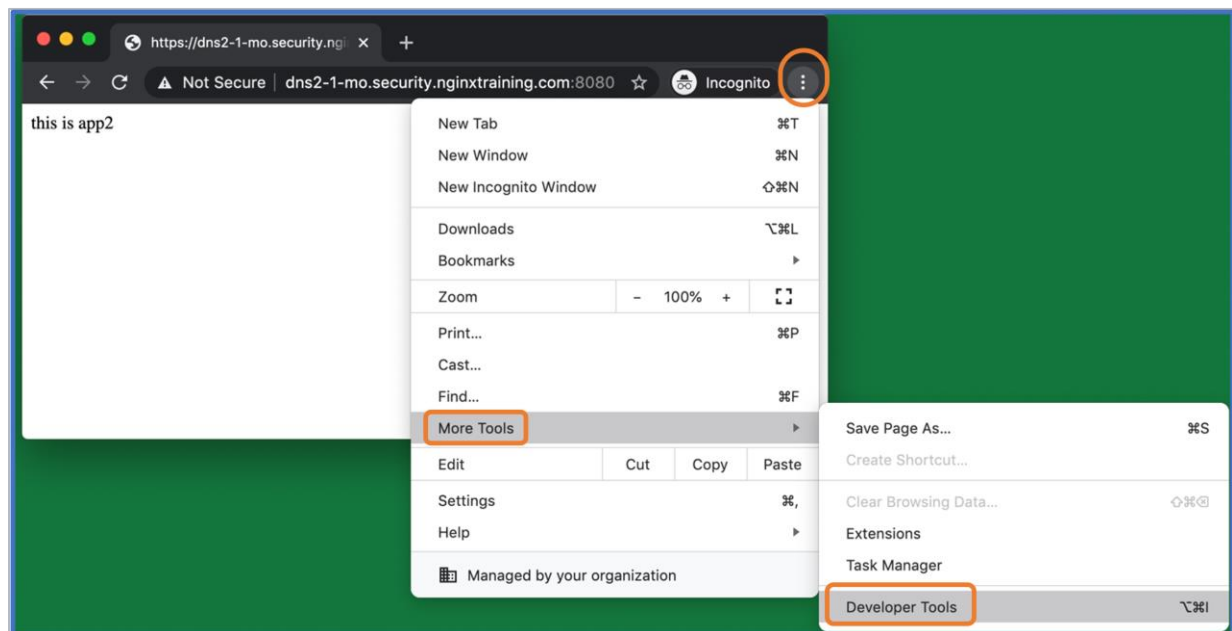
```
https://<www.nginxtraining.com:8080
```



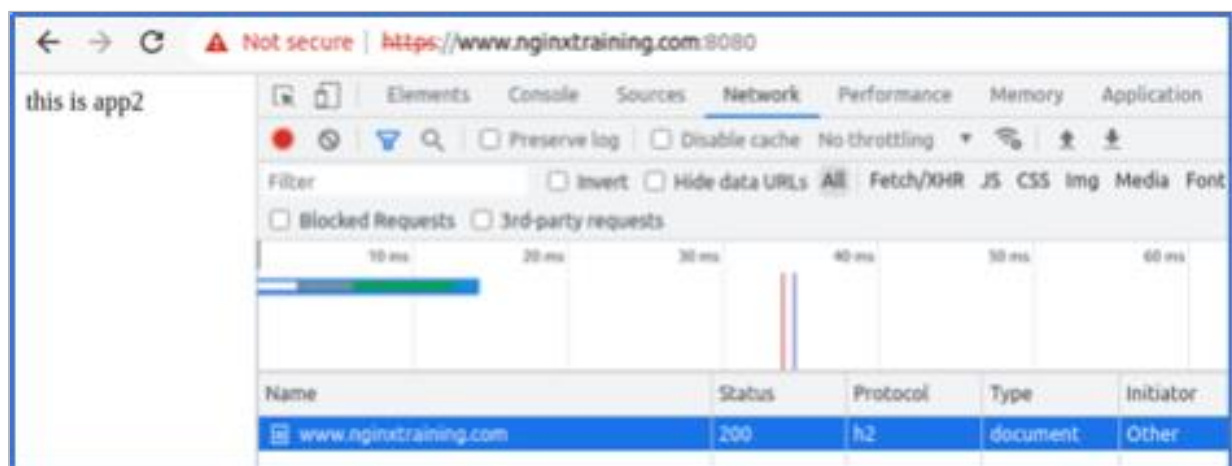
#### NOTE

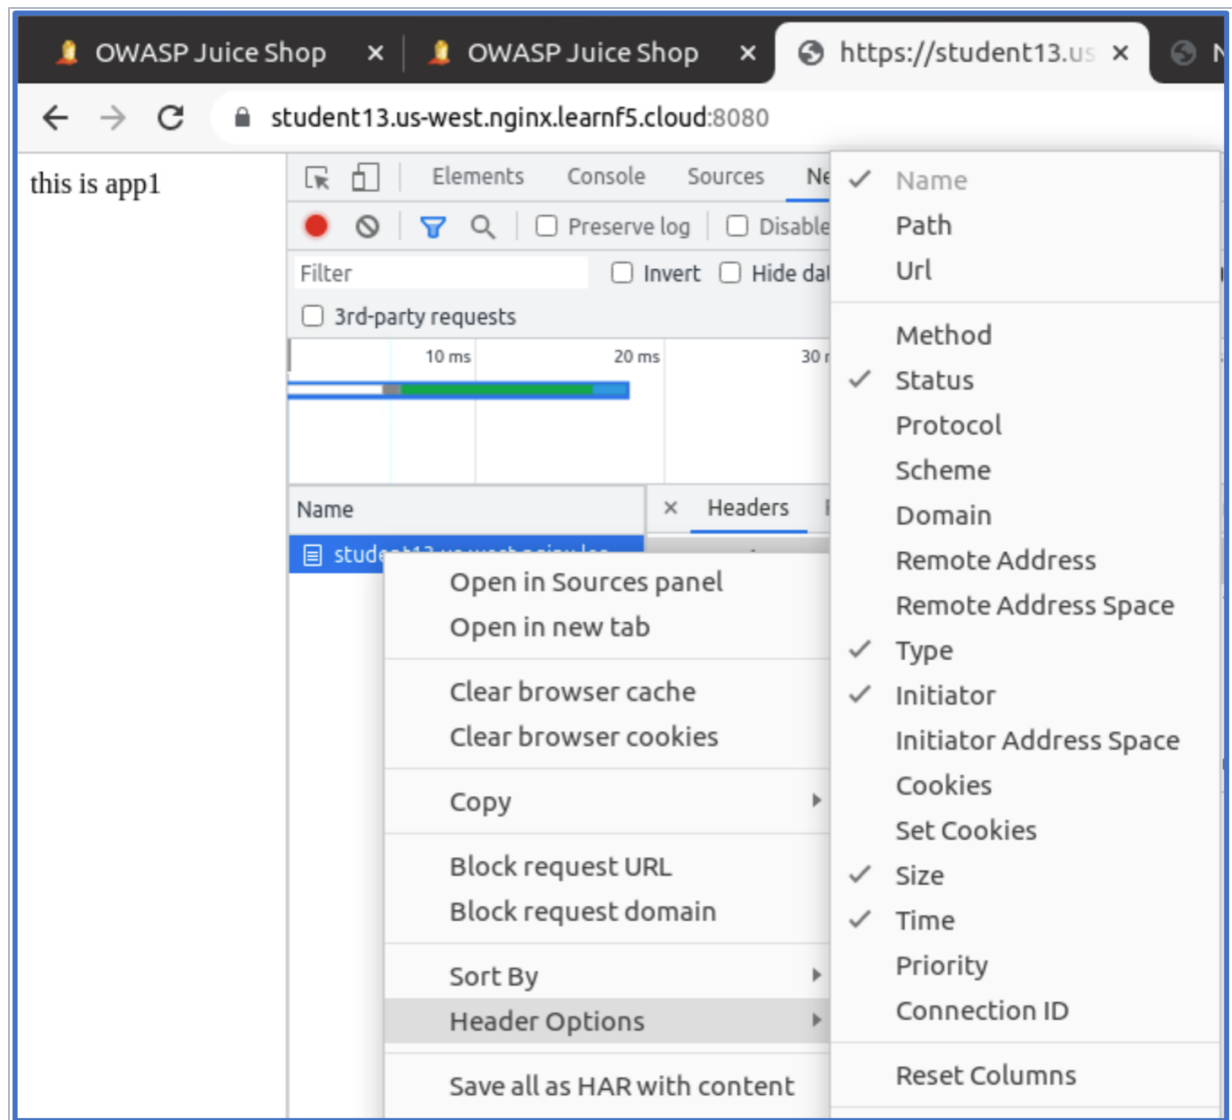
If you use curl to test, it automatically converts HTTP/2 to HTTP1.1 for readability.

20. To verify that your request is using the http2 protocol, you will need to get further information from your browser. For the Chrome browser, click the **3 vertical dots**, select **“More Tools”** and then select **“Developer Tools”**.



21. Click the **“Network”** tab. Refresh the browser web page so that your DNS name is visible. The Status should be 200 and the Protocol **h2** for http2.





#### NOTE



If the Protocol column is not visible, with your lab systems name, [www.nginxtraining.com](https://www.nginxtraining.com), highlighted, right click to select it, choose Header Options and then click on Protocol and finally refresh the JuiceShop page. You might also want to close the console at the bottom of the screen if your systems name is not visible.

### Expected Results

NGINX has now been configured to use the protocol HTTP/2.





## Lab 2: Configure NGINX to not display version number on error pages

Estimated time for completion: **XX minutes**

### Requirements

The following tasks must be completed before beginning this lab:

- The configuration file `/etc/nginx/conf.d/api_server.conf` must exist.

### Objectives

At the end of this lab you will be able to:

- Configure NGINX to not display the version number on error pages.

### Steps

1. In a browser enter your DNS name with port 8080 to display a working webpage. (Close the developer tools.)

`https://www.nginxtraining.com:8080`

2. In a browser enter your systems DNS name with port 8080 followed by the non-existent webpage “shop” so that the default NGINX behavior is displayed showing the version of NGINX.

`https://www.nginxtraining.com:8080/shop`

3. Edit the **nginx.conf** file.

```
$ sudo vim /etc/nginx/nginx.conf
```

4. To not display the NGINX version number on error pages add the directive **server\_tokens off** to the http context. Save the file.

```
http {  
    server_tokens off;  
    include /etc/nginx/mime.types  
    default_type application/octet-stream
```

```
...  
}
```

5. Reload nginx.

```
$ sudo nginx -s reload
```

6. In your browser, test that the NGINX version is no longer displayed when an error page is returned.

```
https://www.nginxtraining.com:8080/shop
```

### Expected Results

NGINX is now configured to not display the version number on error pages.





## Chapter 6

---

### Labs - Secure API Traffic using JWT Authentication

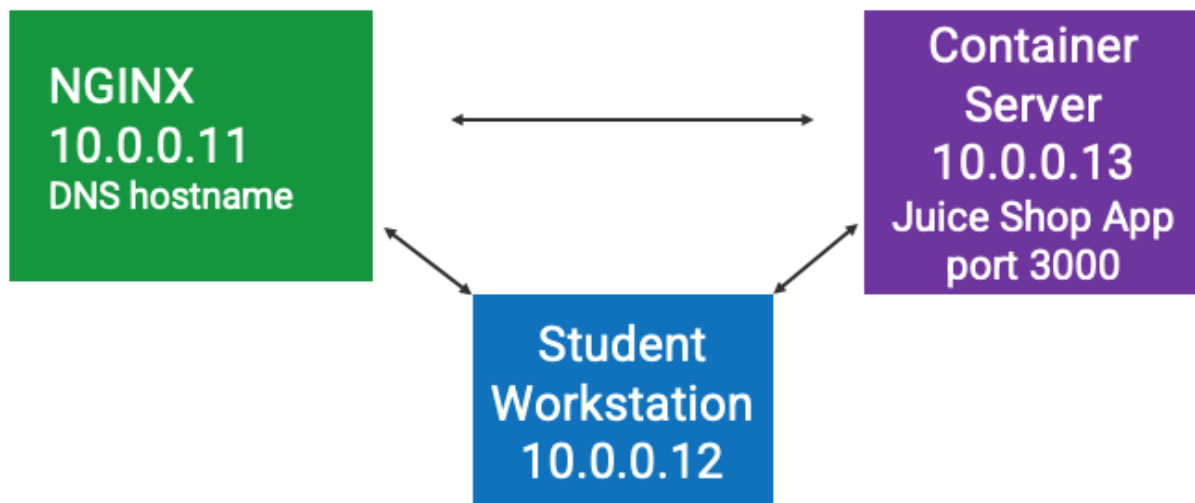
---

#### Lab Contents

1. Create auth\_jwt realm and encode the secret
2. Issue, sign and test the JWT

Estimated time to complete labs: [XX] minutes

#### Lab Network Topology- Possibly Delete Topic?



[Narrative Text]

## Lab 1: Create auth\_jwt Realm and Encode the Secret

Estimated time for completion: **XX minutes**

### Requirements

The following tasks must be completed before beginning this lab:

- The configuration file `/etc/nginx/conf.d/api_server.conf` must exist.

### Objectives

At the end of this lab you will be able to:

- Create an `auth_jwt` realm and encode a secret.
- Issue, sign and test the JWT

### Steps

7. Edit the **api\_server.conf** file.

```
$ sudo vim /etc/nginx/conf.d/api_server.conf
```

8. In the 8080 server context, create the **/products** location with the following **auth\_jwt** and **auth\_jwt\_key\_file** values. Save the file.

```
location /products {  
    auth_jwt "Products API" token=$arg_apijwt;  
    auth_jwt_key_file /etc/nginx/ssl-configs/api_secret.jwk;  
    try_files $uri $uri.html /products.html;  
}
```

**NOTE**

The output should be:



```
server {
    listen 8080 ssl;
    status_zone api_gateway;

    root /usr/share/nginx/html;
    error_log /var/log/nginx/api_server.error.log debug;
    access_log /var/log/nginx/api_server.access.log combined;
    include /etc/nginx/ssl-configs/ssl-params.conf;

    location /products {
        auth_jwt "Products API" token=$arg_apijwt;
        auth_jwt_key_file /etc/nginx/ssl-configs/api_secret.j
        try_files $uri $uri.html /products.html;
    }

    location /api {
        api;
        access_log off;
    }
}
```

9. Use **base64** to create a JSON Web Key (symmetric key used for signing or a secret) using the secret “**fantasticjwt**”. The **tr** command makes the required character substitutions for Base64URL encoding. Copy the command from the **SNIPPETS** file.

```
$ echo -n fantasticjwt | base64 | tr '+/' '-_' | tr -d '='
```

The output should be: **ZmFudGFzdG1jand0**

10. Create the file **api\_secret.jwk** with the JSON Web Key (secret) to be used for signing.

```
$ echo '{"keys":[{"k":"<result from encoded secret>","kty":"oct","kid":"0001"}]}' | sudo tee /etc/nginx/ssl-configs/api_secret.jwk
```

**NOTE**

If you use “**fantasticjwt**” as your secret, you can just copy the command from the **SNIPPETS** file. Otherwise, you must replace the encoded secret with your value from the previous lab step.

Your symmetric key contains these 3 values:

```
{
  "keys": [
    {
      "k": "ZmFudGFzd0ljand0",
      "kty": "oct",
      "kid": "0001"
    }
  ]
}
```

The “**k**” field contains a **base-64 value** of the *<encoded secret>* from the output of the previous step.

The “**kty**” field contains **oct** for the octet sequence.

The “**kid**” field contains **0001** for the serial number.

11. To view the information in a nicer format:

```
$ cat /etc/nginx/ssl-configs/api_secret.jwk | jq
```

12. Test your NGINX configuration/syntax.

```
$ sudo nginx -t
```

13. Use **https://jwt.io** to create the JWT header, payload and signature. Scroll down to get to the **Decoded** section of the webpage.

<https://jwt.io>

14. To create the JWT **HEADER** copy the code from the SNIPPETS file and enter into the HEADER section on the right side of the webpage under **Decoded** replacing the existing header information.

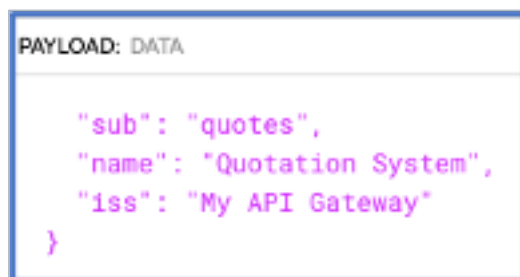




The image shows a web interface titled "Decoded" with a subtitle "EDIT THE PAYLOAD AND SECRET". Below the title is a section labeled "HEADER: ALGORITHM & TOKEN TYPE". Inside this section, a JSON object is displayed in red text: 

```
{  "alg": "HS256",  "typ": "JWT",  "kid": "0001"}
```

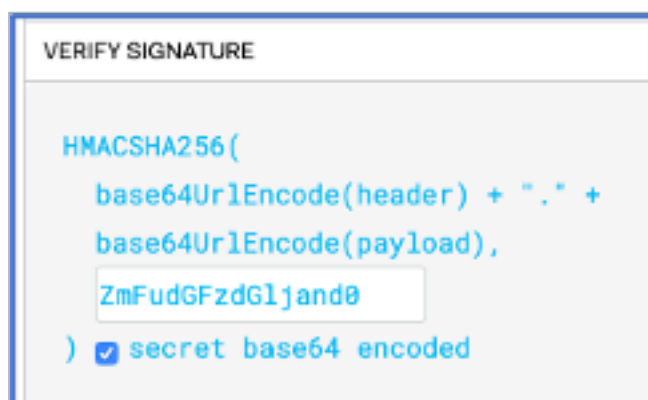
15. To create the JWT **PAYLOAD** copy the code from the SNIPPETS file and enter into the PAYLOAD section replacing the existing information.



The image shows a web interface titled "PAYLOAD: DATA". Inside this section, a JSON object is displayed in purple text: 

```
{  "sub": "quotes",  "name": "Quotation System",  "iss": "My API Gateway"}
```

16. To create the JWT **SIGNATURE** copy the code from the SNIPPETS file or from your previous lab step where you encoded the secret "fantasticjwt" and enter it into the *Verify SIGNATURE* section. Make sure to select **secret base64 encoded**.



The image shows a web interface titled "VERIFY SIGNATURE". Inside this section, there is a code editor with the following text: 

```
HMACSHA256(  base64UrlEncode(header) + "." +  base64UrlEncode(payload),  ZmFudGFzdG1jand0  )
```

 Below the code editor, there is a checkbox labeled "secret base64 encoded" which is checked.

17. On the left side in the **Encoded** section you should have your base64 encoded **HEADER.PAYLOAD.SIGNATURE**. Copy this to use in testing access to your system.

**NOTE**

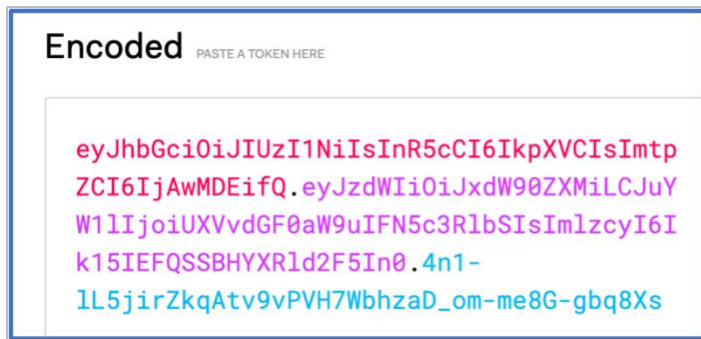
Below the Encoded box, the signature is now verified.



 **Signature Verified**

18. Create the file **quotes.jwt** with the encoded **HEADER.PAYLOAD.SIGNATURE** that you created. You can copy from the SNIPPETS file or from your jwt.io browser.

```
$ sudo vim /etc/nginx/ssl-configs/quotes.jwt
```

**NOTE**

The HEADER is in red, the PAYLOAD in magenta and the SIGNATURE in turquoise.



19. Create the file **products.html** in `/usr/share/nginx/html`.

```
$ sudo echo "<h1>JWT Token WORKS</h1><p>Your JWT Authentication is  
successful</p>" | sudo tee /usr/share/nginx/html/products.html
```

20. Reload NGINX.

```
$ sudo nginx -s reload
```

21. Use the curl command to request /products.

```
$ curl -k https://www.nginxtraining.com:8080/products
```

**NOTE**

You should receive a **401 Authorization Required** response because you set /products to. now require a JWT.

22. Test by appending the file quotes.jwt to the URI argument apijwt.

```
$ curl -k https://www.nginxtraining.com:8080/products?apijwt=`cat /etc/nginx/ssl-configs/quotes.jwt`
```

**NOTE**

The output should be:

```
<h1>JWT Token Works</h1><p>Your JWT Authentication is successful</p>
```

## Expected Results

NGINX has been configured to use a JWT for authentication in order to access a particular URI.



